

# Magnetic-moment-constrained collisions

## Equations, numerical design and verification notes

Project: `sato-morrison`

7 October 2026

### Abstract

These notes derive and audit a small fixed-field JAX prototype of magnetic-moment-constrained collisions. Implemented paths include the local linear weak operator, a positive nodal nonlinear entropy step, five analytic vacuum fields, tangent toroidal Hamiltonian transport, and an explicit finite-range surrogate. Independent homogeneous controls include Lorentz and Dougherty evolution, physical three-velocity Landau weak moments, and direct screened binary encounters. A checked finite-range Fourier calculation shows how spatial localization creates a density nullspace and how Gaussian range lifts it with quartic long-wave decay. Geometric identities, finite-grid conservation and the selected toroidal observable refinements are verified; other velocity-tail convergence and physical collision-rate calibration remain separate requirements. No calibrated dipole confinement, universal zero-set prescription, full Landau kinetic evolution, or publication priority is claimed.

## 1 Scope and equation map

The primary source is Ref. [1]. Its Eq. (35) gives a one-species noncanonical collision structure. Equations (111) and (119) are different levels of description: the former evolves a ten-coordinate pair distribution, while the latter closes two five-coordinate species distributions after a pair-factorization assumption. Equation (181) is a localized linear approximation with a prescribed scalar collision coefficient. This prototype must identify which of these models it implements.

The starting model has a prescribed, nonzero vacuum magnetic field, no magnetic perturbation, a fixed electrostatic reference potential, and one species. The nonlinear local-kernel extension is named separately from literal Eq. (181). No absolute collision rate or laboratory confinement result is established by these calculations.

| Source equations     | Implementation role                                                  |
|----------------------|----------------------------------------------------------------------|
| SM25 (35), (39)–(59) | Weak form, boundary conditions, invariants and stationary family.    |
| SM25 (73)–(100)      | Guiding-center bracket, vacuum assumption and invariant measure.     |
| SM25 (118)–(119)     | Pair closure and multispecies kernel; later stage.                   |
| SM25 (139)–(150)     | Constrained stationary states and initial magnetic-moment marginal.  |
| SM25 (164)–(183)     | Linearization, kernel approximations and the first numerical target. |
| BS25, Sec. V         | Structural comparison to a guiding-center Landau bracket.            |

## 2 Coordinates, measures and the bracket

We use SI variables before nondimensionalization. Let  $u$  be parallel velocity,  $\mu = mv_\perp^2/(2B)$  the magnetic moment, and  $\eta = \mu B$  the perpendicular kinetic energy. With  $\mathbf{b} = \mathbf{B}/B$ ,

$$\mathbf{B}^* = \mathbf{B} + \frac{mu}{q} \nabla \times \mathbf{b}, \quad B_\parallel^* = \mathbf{b} \cdot \mathbf{B}^*. \quad (1)$$

For the vacuum backgrounds used in Ref. [1],  $\mathbf{b} \cdot \nabla \times \mathbf{b} = 0$ , hence  $B_\parallel^* = B$ . This does not imply  $\nabla \times \mathbf{b} = 0$ : curvature terms in  $\mathbf{B}^*$  generally remain.

In coordinates  $z = (\mathbf{X}, u, \mu)$  the gyrophase-independent bracket acts on a covector  $a = (\mathbf{a}_X, a_u, a_\mu)$  as

$$\mathcal{J}_\mu a = \begin{pmatrix} (\mathbf{b} \times \mathbf{a}_X)/(qB) + \mathbf{B}^* a_u/(mB) \\ -\mathbf{B}^* \cdot \mathbf{a}_X/(mB) \\ 0 \end{pmatrix}. \quad (2)$$

The Hamiltonian and characteristic equations are

$$E = \frac{1}{2}mu^2 + \mu B + q\Phi, \quad (3)$$

$$\dot{\mathbf{X}} = u \frac{\mathbf{B}^*}{B} + \frac{\mathbf{b} \times (\mu \nabla B + q \nabla \Phi)}{qB}, \quad (4)$$

$$\dot{u} = -\frac{\mathbf{B}^* \cdot (\mu \nabla B + q \nabla \Phi)}{mB}, \quad \dot{\mu} = 0. \quad (5)$$

The signs in (2) must be checked against the cross-field drift in (4), not inferred from a memorized matrix convention.

The phase-space measures are

$$d\Gamma = B d^3X du d\mu = d^3X du d\eta. \quad (6)$$

We follow the article's normalization:  $n(\mathbf{X}) = \int f du d\eta$ . If  $F$  is an ordinary gyrotropic velocity density normalized with  $d^3v$ , then  $d^3v = (2\pi B/m)dud\mu$  and  $f = (2\pi/m)F$ . Track this constant, particularly in multispecies comparisons.

The coordinate change  $\zeta = (\mathbf{X}, u, \eta)$  has derivative

$$T = \frac{\partial \zeta}{\partial z} = \begin{pmatrix} I_3 & 0 & 0 \\ 0 & 1 & 0 \\ \mu(\nabla B)^T & 0 & B \end{pmatrix}, \quad \mathcal{J}_\eta = T \mathcal{J}_\mu T^T. \quad (7)$$

This constructs the matrix in SM25 (93), equivalently (182), without duplicating its entries by hand. It also gives

$$\mathcal{J}_\eta \nabla_\zeta (\eta/B) = 0, \quad \dot{\eta} = \mu \dot{\mathbf{X}} \cdot \nabla B. \quad (8)$$

Thus a discretization that freezes  $\eta$  during spatial transport is not the magnetic-moment-preserving model in a nonuniform field.

With general coordinates and measure  $\rho(z)dz$ , conservative divergence means  $\rho^{-1}\partial_a(\rho F^a)$ , not  $\partial_a F^a$ . Verify the Liouville identity  $\partial_a(\rho \mathcal{J}^{ab}) = 0$ . The prototype uses  $\rho = B$  in the  $\mu$  chart and  $\rho = 1$  in the article's  $\eta$  chart.

## Continuous geometric audit

The implemented analytic fields pass automatic-differentiation checks of  $\nabla \cdot \mathbf{B} = 0$ ,  $\nabla \times \mathbf{B} = 0$ , antisymmetry, the magnetic-moment Casimir, the Liouville density, and the Jacobi identity in both charts, at independent off-axis points and for both charge signs. Centered differences of  $\mathbf{B}$  show second-order convergence to the automatic derivatives for the rational fields. These sampled identities audit the continuous formulas; they do not prove a discrete Jacobi identity for the collision mesh or positivity of  $B$  throughout an unsampled domain. Host validation rejects unsupported fields, nonfinite inputs, zero field, excluded dipole/toroidal sources, nonpositive mass, zero charge and negative moment.

For the vacuum toroidal field  $B = C/R$ , the characteristic equations give  $\dot{R} = \dot{u} = \dot{\mu} = 0$ ,  $\dot{\theta} = u/R$  and  $\dot{z} = (mu^2 + \mu B)/(qC)$ . Thus annular radial boundaries are tangent to Hamiltonian flow; a periodic  $z$  and toroidal angle are compatible with this geometry. The library checks this identity and the uniform electrostatic drift independently of collision evolution.

## Normalization is part of the closure

Choose fixed scales  $L_*, v_*, E_*, t_*$  and a diagonal coordinate scale matrix  $S$ . For  $\zeta = S\hat{\zeta}$ ,  $E = E_*\hat{E}$  and  $t = t_*\hat{t}$ ,

$$\hat{\mathcal{J}} = t_* E_* S^{-1} \mathcal{J} S^{-T}. \quad (9)$$

State these scales with every dimensional result. A Euclidean phase-space projector applied to a vector containing spatial velocities, accelerations and energy rates is not defined independently of those scales.

For a purely affine chart change  $y = Az$  with time and energy conventions fixed, a pair vector transforms as  $d_y = Ad_z$ . Its quadratic form is unchanged only if

$$\Pi_y = A^{-T} \Pi_z A^{-1}, \quad d_y^T \Pi_y e_y = d_z^T \Pi_z e_z. \quad (10)$$

Recreating an identity matrix and Euclidean projector in the new chart generally specifies a different kernel. For the position-dependent map (7), compute the two vectors in one common chart before subtracting:

$$\Delta_\eta \phi = T(z) \mathcal{J}_\mu(z) \nabla_z \phi - T(z') \mathcal{J}_\mu(z') \nabla_{z'} \phi'. \quad (11)$$

This is the safe route when storing the state on  $\mu$  nodes but reproducing a kernel defined in the  $\eta$  chart.

## 3 Nonlinear collision structure and conservation

For a single species define

$$\Delta \phi = \mathcal{J} \nabla \phi - \mathcal{J}' \nabla' \phi', \quad \xi = \Delta E. \quad (12)$$

A specified kernel must obey exchange symmetry, symmetry of its quadratic form, positive semidefiniteness and the energy degeneracy:

$$\Pi = \Pi^T \succeq 0, \quad \Pi \xi = 0. \quad (13)$$

The source strong form is

$$C[f] = \nabla_z \cdot \left[ f \mathcal{J} \int f' \Pi (\mathcal{J}' \nabla' \ln f' - \mathcal{J} \nabla \ln f) dz' \right] \quad (14)$$

in a flat invariant-measure chart. Integration by parts, antisymmetry of  $\mathcal{J}$  and exchange of the two labels give

$$\int \phi C[f] d\Gamma = -\frac{1}{2} \iint f f' \Delta \phi^T \Pi \Delta \ln f d\Gamma d\Gamma'. \quad (15)$$

All following equalities require the boundary terms to vanish. The paper states conditions for both the total flux and the Hamiltonian part; a finite numerical box does not satisfy them automatically.

Set  $\phi = 1$  in (15) to obtain particle conservation. Set  $\phi = E$  and use (13) for collision energy conservation. Since  $\mathcal{J}\nabla g(\mu) = 0$  for every smooth  $g$ ,

$$\frac{d}{dt} \int f g(\mu) d\Gamma = 0. \quad (16)$$

For closed evolution this is equivalent to conserving the full marginal

$$G(\mu) = \int B f(\mathbf{X}, u, \mu) d^3 X du. \quad (17)$$

It is stronger than conservation of  $\int \mu f d\Gamma$  alone.

Define  $S[f] = - \int f \ln(f/f_{\text{ref}}) d\Gamma$  with a constant reference scale inside the logarithm. Particle conservation removes that scale from the entropy rate. Equation (15) yields

$$\dot{S} = \frac{1}{2} \iint f f' \Delta \ln f^T \Pi \Delta \ln f d\Gamma d\Gamma' \geq 0. \quad (18)$$

These are structural statements from the model, not a proof of a unique final state. Additional null directions of  $\Pi$  can prevent mixing.

A candidate observable  $p$  is conserved by the full kinetic model only if the Hamiltonian flow preserves it and the collision form annihilates  $\Delta p$ . An external dipole field does not imply conservation of all Cartesian momenta. For an axisymmetric model, check canonical angular momentum using the actual kernel and boundaries rather than inferring it from the field name.

## 4 The fixed-field linear model

Write  $f = f_0(1 + h)$ , where  $f_0 > 0$  is stationary. For the source stationary family,

$$f_0 = A \exp[-\beta E_0 + g(\mu)], \quad \mathcal{J}\nabla \ln f_0 = -\beta \mathcal{J}\nabla E_0. \quad (19)$$

With the fields and kernel fixed, the zeroth-order collision term vanishes. This supplied energy/moment family is not claimed to exhaust stationary states. For the exactly local poloidal surrogate, (68) adds an arbitrary conserved-flux function. The first-order weak form is

$$\int \phi C_L[f_0 h] d\Gamma = -\frac{1}{2} \iint f_0 f'_0 \Delta \phi^T \Pi_0 \Delta h d\Gamma d\Gamma'. \quad (20)$$

The local scalar approximation in SM25 (181) uses

$$\Pi_0 = D P I_x P \delta^3(\mathbf{X} - \mathbf{X}'), \quad P = I - \frac{\xi_0 \xi_0^T}{\xi_0^T \xi_0}, \quad I_x = \text{diag}(1, 1, 1, 0, 0). \quad (21)$$

Here  $D$  is a prescribed constant, and the projector is evaluated in a specified normalized common chart. The two velocity integrals remain after the spatial delta function is used. In a  $\mu$  chart, the local pair measure contains two velocity Jacobians; do not accidentally retain only one factor of  $B$ .

The perturbation norm

$$Q[h] = \frac{1}{2} \int f_0 h^2 d\Gamma \quad (22)$$

satisfies

$$\dot{Q} = -\frac{1}{2} \iint f_0 f'_0 \Delta h^T \Pi_0 \Delta h \, d\Gamma d\Gamma' \leq 0 \quad (23)$$

for collision-only dynamics. A compatible stationary Hamiltonian flow also preserves  $Q$  under suitable boundary conditions.

For self-consistent fields, (20) is not a complete linearization unless the omitted terms have been ordered out. Differentiating the exact energy-degeneracy condition gives

$$\delta \Pi \xi_0 + \Pi_0 \delta \xi = 0. \quad (24)$$

The approximations in SM25 (164)–(170) must be carried explicitly into any extension. The fixed-field implementation checks its linear operator against a directional derivative of the nonlinear residual; this does not validate omitted self-consistent-field terms.

## 5 Uniform-field reduction and exact null modes

This section is an independent derivation from the source Eq. (181). Take a uniform field, a constant reference potential and  $f_0 = f_0(u, \eta)$ . For  $\mathbf{b} = \mathbf{e}_z$ ,

$$\mathcal{J}a = \left( -\frac{a_y}{qB}, \frac{a_x}{qB}, \frac{a_u}{m}, -\frac{a_z}{m}, 0 \right)^T, \quad \mathcal{J}\nabla E_0 = (0, 0, u, 0, 0)^T. \quad (25)$$

Consequently  $\xi_0 = (u - u')\mathbf{e}_z$  in the spatial block, and

$$PI_x P = \text{diag}(1, 1, 0, 0, 0). \quad (26)$$

At  $u = u'$  use the continuous limit (26); this prescription does not extend automatically to a general nonuniform field.

Only perpendicular spatial derivatives survive. Let

$$n_0 = \int f_0 \, du \, d\eta, \quad \delta n = \int \delta f \, du \, d\eta. \quad (27)$$

Two applications of the cross-product matrix give  $(\mathbf{b} \times)^2 = -I_\perp$ . Substitution into the strong form gives

$$\boxed{C_L[\delta f] = \frac{D}{(qB)^2} \nabla_\perp^2 (n_0 \delta f - f_0 \delta n).} \quad (28)$$

This is spatial mixing of velocity-dependent structure, not ordinary homogeneous velocity-space thermalization.

For a Fourier mode with  $\mathbf{k}_\perp = \mathbf{k} - (\mathbf{k} \cdot \mathbf{b})\mathbf{b}$ , define

$$\lambda = \frac{D n_0 |\mathbf{k}_\perp|^2}{(qB)^2}, \quad \bar{h} = \frac{1}{n_0} \int f_0 h \, du \, d\eta. \quad (29)$$

The collision-only equation is

$$\partial_t h = -\lambda(h - \bar{h}), \quad h(t) = \bar{h} + [h(0) - \bar{h}]e^{-\lambda t}. \quad (30)$$

Thus a density-like perturbation is undamped. If  $\mathbf{k}_\perp = 0$ , every velocity perturbation is collisionally undamped. A homogeneous anisotropic distribution is therefore not an isotropization benchmark for this model.

With positive velocity quadrature weights  $w_i$ , set

$$m_i = w_i f_{0i}, \quad s_i = \sqrt{m_i}, \quad n_0 = \sum_i m_i. \quad (31)$$

The matrix for  $h$  and its entropy-scaled form are

$$L_h = -\lambda \left( I - \frac{\mathbf{1}\mathbf{m}^T}{n_0} \right), \quad (32)$$

$$M_v^{1/2} L_h M_v^{-1/2} = -\lambda \left( I - \frac{\mathbf{s}\mathbf{s}^T}{n_0} \right). \quad (33)$$

For  $\lambda > 0$  the spectrum has one zero eigenvalue and  $N_v - 1$  copies of  $-\lambda$ . This is an exact finite-quadrature statement when the same  $n_0$  is used throughout. It provides a stringent factor, sign and normalization test.

When parallel streaming is added,  $\partial_t h = -ik_{\parallel} u h + L_h h$ . A density-like collision null mode need not remain in that nullspace under streaming. Distinguish the collision spectrum from the spectrum of the combined generator. With  $\mathbf{k}_{\perp} = 0$  and frozen fields the exact streaming solution is  $h(t, u, \eta) = h(0, u, \eta) e^{-ik_{\parallel} u t}$ ; this is not self-consistent Landau damping. An executed distribution regression integrates the Fourier generator obtained from the actual uniform Poisson characteristic with an independent DOP853 integrator and compares it with this analytic phase at four times through  $t = 1.1$ . It uses  $m = 1.7$ ,  $B = 2.3$ , charges  $-0.8$  and  $1.3$ , and  $k_{\parallel} = 2$ . The positive reconstructed distribution agrees within the stated  $2 \times 10^{-12}$  relative/ $3 \times 10^{-13}$  absolute test tolerances; each velocity-bin spatial population is preserved. The independent pair oracle gives zero parallel collision action at nonzero  $D$ . A separate active perpendicular-grid regression sets  $D = 0$  and verifies that the implicit collision step is the identity. These finite-time controls do not assert long-time damping or resolve a kinetic recurrence limit.

## 6 A finite-range surrogate and its exact density branch

This section specifies a new calculation for an explicit surrogate, not a Coulomb scattering tensor or a claim of publication priority. Retain the uniform field and constant-potential assumptions of (25), and replace the spatial delta function in (21) by a nonnegative, even, translation invariant kernel  $W(\mathbf{X} - \mathbf{X}')$  with  $\int W d^3X = 1$ . Use the same uniform  $PI_x P = I_{\perp}$  prescription, including its continuous equal- $u$  limit. The weak form then gives

$$C_W[\delta f] = \frac{D}{(qB)^2} \nabla_{\perp}^2 [n_0 \delta f - f_0(W * \delta n)]. \quad (34)$$

To check the sign and convolution placement, the projected pair difference for  $h$  is  $(\mathbf{b} \times)/(qB)$  applied to  $\nabla_{\perp} h(\mathbf{X}, v) - \nabla'_{\perp} h(\mathbf{X}', v')$ . The first self term integrates to  $n_0 \nabla_{\perp} h$ ; the cross term is  $\nabla_{\perp} W * \delta n$ . The second cross-product operation contributes  $-I_{\perp}$ , giving (34) with the displayed diffusion sign.

**Finite-quadrature proposition.** Let  $m_i = w_i f_{0i} > 0$ ,  $n_0 = \sum_i m_i$ ,  $s_i = \sqrt{m_i}$ ,  $r(\mathbf{k}) = \widehat{W}(\mathbf{k})$ , and  $\lambda = D n_0 |\mathbf{k}_{\perp}|^2 / (qB)^2$ . In the entropy coordinates,

$$-M_v^{1/2} L_{W, \mathbf{k}} M_v^{-1/2} = \lambda \left[ I - r(\mathbf{k}) \frac{\mathbf{s}\mathbf{s}^T}{n_0} \right]. \quad (35)$$

Its exact decay rates are  $\lambda$  on the  $N_v - 1$  dimensional velocity-neutral subspace and  $\lambda[1 - r(\mathbf{k})]$  on the density direction  $\mathbf{s}$ . Indeed,  $\mathbf{s}^T \mathbf{s} = n_0$ , so the normalized outer product is an orthogonal rank-one

projector. An even nonnegative kernel has real  $r \leq 1$ ; thus all rates are nonnegative. This proof does not require  $r \geq 0$ . For a discrete periodic spatial quadrature the identical statement holds with  $r$  equal to the Fourier multiplier of the sampled, symmetrically normalized kernel, rather than the continuum multiplier.

For a normalized periodized isotropic Gaussian of covariance  $\ell^2 I$ ,  $r = e^{-\ell^2 |\mathbf{k}|^2/2}$  and hence

$$\lambda_{\text{density}} = \frac{Dn_0 |\mathbf{k}_\perp|^2}{(qB)^2} \left(1 - e^{-\ell^2 |\mathbf{k}|^2/2}\right) = \frac{Dn_0 \ell^2}{2(qB)^2} |\mathbf{k}_\perp|^2 |\mathbf{k}|^2 + O(|\mathbf{k}_\perp|^2 \ell^4 |\mathbf{k}|^4). \quad (36)$$

A transverse long wave therefore has quartic rather than quadratic density decay. Every finite nonzero Gaussian range lifts the nonzero-transverse-wave local density null mode, but the lifted rate vanishes as  $\ell^2$  in the local limit. Homogeneous and purely parallel modes remain fully collisionally null. Consequently this calculation does not establish a whole-space spectral gap, or a gap for all modes on a periodic domain containing parallel-only modes. On a fixed finite set of transverse Fourier modes, it establishes the stated positive density rates. The limits  $t \rightarrow \infty$  and  $\ell \rightarrow 0$  need not commute for the density branch.

The independent NumPy check in `examples/06_range.py` assembles the ordered spatial/velocity pair weak sum, including the factor  $1/2$ . The Fourier basis is differentiated exactly; the Gaussian is evaluated by positive periodic image sums and normalized with one scalar to retain pair symmetry. Three spatial resolutions, three velocity node counts and three kernel ranges are compared with (35). The velocity weights are normalized once at setup to fixed  $n_0$ ; this verifies the finite-quadrature theorem rather than physical velocity-tail convergence. The result is independent of the production JAX pair implementation.

## Literature comparison and novelty limits

SM25 supplies the noncanonical weak structure and its local scalar approximation. SM24 supplies the earlier noncanonical construction [2]. BS25, Eqs. (99)–(103), retains magnetic-moment friction and diffusion through its gyrophase-dependent projection  $\Delta_{gc}^\mu$ , and explains the structural reduction obtained by removing that component and replacing the relative velocity by the reduced Hamiltonian pair velocity. That passage does not identify every scalar/local SM25 approximation with the complete conventional Landau tensor. Our uniform surrogate remains unable to isotropize a homogeneous distribution.

KH17 supplies the discrete-gradient conservation and entropy method; those properties do not supply a coercivity constant. Bressan, Kraus, Maj and Morrison [3] analyze sufficient relaxation conditions, including additional finite-dimensional inequalities and tentative infinite-dimensional extensions. Their framework motivates testing the actual nullspace before assigning asymptotic decay to a constrained maximum-entropy state.

On 7 October 2026, primary-source searches of the SM25 title combined with finite range, nullspace and Gaussian terms, and searches for noncanonical finite-range collisions and magnetic-moment quartic relaxation, found no matching statement of (35)–(36). Local full-text checks covered SM25, SM24, BS25, KH17 and BBQ15. This is a scoped independently checked derivation for the specified surrogate; absence from these searches is not proof of publication priority. No general nonuniform finite-range theorem, Coulomb rate, or physical metastability is inferred.

## 7 Stationary states and initial constraints

The source stationary family (19) follows from energy degeneracy and the magnetic-moment Casimir. For  $g(\mu) = -\beta\gamma\mu$ ,

$$f_* = A \exp \left[ -\beta \left( \frac{1}{2} m u^2 + \mu(B + \gamma) + q\Phi \right) \right], \quad B + \gamma > 0. \quad (37)$$

Integrating over  $u$  and  $\mu$  using (6) gives

$$n_*(\mathbf{X}) = A \sqrt{\frac{2\pi}{\beta m}} e^{-\beta q\Phi} \frac{B}{\beta(B + \gamma)}. \quad (38)$$

The corresponding temperature parameters in energy units are

$$T_{\parallel} = \beta^{-1}, \quad T_{\perp} = \frac{B}{\beta(B + \gamma)}. \quad (39)$$

These formulas are useful checks of velocity integration and the definition of pressure. They do not establish that an arbitrary initial state evolves to this particular family member.

Two further velocity-integral regressions check explicit limits of the known stationary family. With  $f = \exp[-\beta(mu^2/2 + \mu(B + \gamma))]$  and the specified measure  $Bdu d\mu$ , its density is  $n(B) = c_{\beta}B/(B + \gamma)$ , where  $c_{\beta} = \sqrt{2\pi}/(\beta m)/\beta$ . At  $\gamma = 0$ , independent positive-quadrature integrals give the constant  $c_{\beta}$  at three points in each of the five fields, including nonunit mass. For the harmonic mirror of strength  $B_0$  and amplitude  $a$ ,

$$B(a) = B_0 + a(z^2 - R^2/2) + O(a^2), \quad n(B(a)) = \frac{c_{\beta}B_0}{B_0 + \gamma} + a \frac{c_{\beta}\gamma(z^2 - R^2/2)}{(B_0 + \gamma)^2} + O(a^2).$$

The quadrature density agrees with the independently integrated analytic formula, and its divided-difference slope approaches this expansion under three amplitude halvings. These tests concern the known density family, not attraction to it, physical rate calibration or the near-uniform projector limit.

For a prescribed fixed-field energy and a known initial marginal  $G_0(\mu)$ , let

$$Z_{\mu}(\beta) = \int B(\mathbf{X}) e^{-\beta E(\mathbf{X}, u, \mu)} d^3 X du. \quad (40)$$

A candidate constrained maximum-entropy state is

$$f_*(\mathbf{X}, u, \mu) = \frac{G_0(\mu)}{Z_{\mu}(\beta)} e^{-\beta E(\mathbf{X}, u, \mu)}. \quad (41)$$

Solve for  $\beta$  from total energy. Add other invariant multipliers if they actually survive the chosen dynamics. The same initial inverse-temperature parameter need not equal the final multiplier. Extra kernel invariants can make (41) inaccessible.

At fixed  $G_0$ , differentiation of the mean energy of (41) gives

$$\frac{dE_*}{d\beta} = - \int G_0(\mu) \text{Var}_{\beta, \mu}(E) d\mu \leq 0. \quad (42)$$

It is strictly negative if at least one populated moment bin has nonconstant energy. Thus an energy target within the attainable bracket selects a unique  $\beta$  in that bracket. Positive-temperature existence is not automatic for an arbitrary truncated initial distribution; the implementation rejects an unbracketed

target. The same calculation holds for positive finite quadrature. With matching energy and full moment marginal, relative entropy gives  $S[f_*] - S[f] = \int f \ln(f/f_*) d\Gamma \geq 0$ . This identifies a constrained maximum-entropy candidate; it does not establish dynamic accessibility under an operator possessing additional invariants.

For self-consistent electrostatics,

$$-\epsilon_0 \nabla^2 \Phi = \sum_s q_s n_s + \rho_{\text{fixed}}. \quad (43)$$

The self-field energy is  $\epsilon_0 \int |\nabla \Phi|^2 d^3 X / 2$ , with boundary terms treated consistently. A prescribed external potential contributes  $\sum_s \int q_s \Phi_{\text{ext}} f_s d\Gamma_s$  without the self-interaction half-factor. A periodic field solve requires net neutrality and a gauge choice.

## 8 Comparison models and what each tests

The conventional reduced Landau literature retains magnetic-moment friction and diffusion even when the distribution is gyrophase independent [4, 5, 6, 7]. Brizard and Sugama explicitly discuss a structural reduction with  $\Delta_{\text{gc}}^\mu = 0$  and a reduced relative velocity in Sec. V of Ref. [7]. This is a bridge for comparison, not a verification that every simplified Sato–Morrison kernel is identical to their Landau tensor. The papers use different conventions; convert all masses, charges and unit factors before comparing implementations.

### Pitch-angle scattering

With  $v = |\mathbf{v}|$  and  $\xi = u/v$ , use

$$C_{\text{Lz}}[F] = \frac{\nu(v)}{2} \partial_\xi [(1 - \xi^2) \partial_\xi F]. \quad (44)$$

The Legendre identity gives  $C_{\text{Lz}}[P_\ell] = -\nu \ell(\ell + 1) P_\ell / 2$  for constant  $\nu$  on a speed shell. Hence degrees 0, 1, 2, 3 have rates 0,  $\nu$ ,  $3\nu$ ,  $6\nu$ . Speed-shell population and energy are conserved, but momentum of the evolved species can be transferred to the fixed reservoir. Since  $\mu = mv^2(1 - \xi^2)/(2B)$ , this is a moment-changing comparison. Derive its coordinate transform; it is not bare diffusion at fixed  $u$  in  $\mu$ .

### A nonlinear conserving Fokker–Planck model

The single-species Dougherty form used as a comparison is

$$C_D[F] = \nu \nabla_v \cdot [(\mathbf{v} - \mathbf{U})F + \theta \nabla_v F], \quad \mathbf{U} = \frac{1}{n} \int \mathbf{v} F d^3 v, \quad \theta = \frac{1}{3n} \int |\mathbf{v} - \mathbf{U}|^2 F d^3 v. \quad (45)$$

Integration by parts shows number and momentum conservation. The trace of the centered covariance is constant. Its matrix satisfies

$$\dot{\Theta} = -2\nu(\Theta - \theta I), \quad \Theta(t) = \theta I + [\Theta(0) - \theta I] e^{-2\nu t}. \quad (46)$$

For an initially Gaussian homogeneous state this specifies the full Gaussian solution. Use it as an exact code test. A fixed bath temperature or drift would define a different model. Neither version is the exact Coulomb operator.

## Physical three-velocity Landau reference

A one-species normalized reference is

$$U(\mathbf{w}) = \frac{I - \hat{\mathbf{w}}\hat{\mathbf{w}}^T}{|\mathbf{w}|}, \quad (47)$$

$$C_{\text{La}}[F] = \Gamma \nabla_v \cdot \int F F' U(\mathbf{v} - \mathbf{v}') (\nabla_v \ln F - \nabla_{v'} \ln F') d^3 v'. \quad (48)$$

Derive the dimensional prefactor for the selected normalization separately. Its positive kernel and velocity-difference null direction support conventional number, momentum, energy and entropy tests. Use an independently assembled small 3V problem before implementing a general transformed operator.

A gyrotropic two-coordinate distribution still has three-dimensional collision geometry. Converge the gyrophase integral or derive its analytic average; a literal 2V Landau kernel is not an equivalent comparator. Maxwell-molecule and BKW-type tests in Ref. [8] can test a different kernel's numerical machinery, but do not establish Coulomb rates.

## 9 Analytic magnetic fields and boundaries

A vacuum field can be specified locally as the gradient of a harmonic scalar potential. Use the following fields with explicit domains.

**Uniform field.**  $\mathbf{B} = B_c \mathbf{b}_0$ , with  $|\mathbf{b}_0| = 1$ .

**Harmonic mirror.** Let  $R^2 = x^2 + y^2$  and

$$\Psi_B = B_c z + a \left( z^3/3 - zR^2/2 \right), \quad (49)$$

$$\mathbf{B} = \nabla \Psi_B = (-azx, -azy, B_c + a[z^2 - R^2/2]). \quad (50)$$

Its curl and divergence vanish identically. Choose a domain on which  $B$  stays strictly positive. The dimensionless variation is controlled by  $aL^2/B_c$ . An arbitrary field  $B(z)\mathbf{e}_z$  is not an admissible substitute.

**Toroidal vacuum field.** On  $R > R_{\min} > 0$ ,

$$\mathbf{B} = \frac{C}{R^2}(-y, x, 0). \quad (51)$$

This provides curvature and field-strength variation without a parallel gradient of  $B$ . Its noncontractible domain is part of the problem.

**Point dipole.** Away from the source,

$$\mathbf{B} = \frac{C_d}{r^5}(3xz, 3yz, 2z^2 - x^2 - y^2), \quad \Psi_B = -C_d z/r^3. \quad (52)$$

This is a field test, not a model of a finite winding or material boundary. Exclude  $r = 0$  and any near-source region outside the guiding-center ordering.

**Symmetry breaking.** A controlled addition  $\delta\Psi_B = \varepsilon B_c xy/L$  is harmonic. Its gradient can be added to a vacuum background; check the resulting nonzero-field domain and boundary setup.

Do not make a nonperiodic magnetic field periodic by wrapping its array. Begin with collision-only nonuniform tests under natural zero-collision-flux boundary conditions. For full dynamics, require tangent Hamiltonian flux, a derived invariant-compatible reflection, a specified confining potential, or an open system with measured flux budgets. Arbitrary specular reflection need not preserve  $\mu$ . A density-profile test without these choices is not a closed relaxation experiment.

## Complete smooth nullspace of the local toroidal surrogate

The following classification concerns the specified local  $PI_X P$  closure, not a conventional Coulomb operator. Assume  $B = (C/R)\mathbf{e}_\theta$ ,  $C > 0$ ,  $m > 0$ ,  $q \neq 0$ , constant electrostatic potential and  $D > 0$ . Take a connected radial interval away from  $R = 0$ , periodic  $\theta, z$ , and a product velocity domain  $I_u \times I_\mu$  with both intervals of positive length. The equilibrium weight is strictly positive in this domain. Consider real  $C^2$  test functions in the weighted form domain, so the pair quadratic integral is finite. Periodicity and natural zero collision flux on nonperiodic faces give the weak identity (20).

**Proposition.** The smooth collision-null functions are exactly

$$h = \phi(R, \theta, z) + g(\theta, \mu) + a(\theta)E + b(\theta)mRu, \quad E = \frac{1}{2}mu^2 + \mu C/R. \quad (53)$$

Here the functions are arbitrary subject to the stated regularity, periodicity and integrability. Redundant angle-only terms may be moved between  $\phi$  and  $g$ . The common collision and ideal stationary functions are exactly

$$h = \phi(R) + g(\mu) + aE + bmRu, \quad a, b \text{ constant.} \quad (54)$$

For a stationary positive  $f_0$ , skew Hamiltonian transport and the collision norm identity imply that these also exhaust the smooth zero modes of the combined linear generator. This classifies stationary null modes; it does not prove a decay rate, coercivity constant, or spectral gap.

**Proof of the collision statement.** In the common Cartesian  $\eta$  chart, write the spatial action in the local orthonormal basis ( $\mathbf{e}_R, \mathbf{e}_\theta, \mathbf{e}_z$ ):

$$A_X h = \left( \frac{Rh_z}{qC}, \frac{h_u}{m}, \frac{-Rh_R + uh_u}{qC} \right), \quad V_X = A_X E = \left( 0, u, \frac{mu^2 + \mu C/R}{qC} \right). \quad (55)$$

The factor  $m$  in  $\mathbf{B}^*$  cancels the bracket denominator  $m$  in the curvature term, so the last numerator contains  $uh_u$ , not  $muh_u$ . The energy characteristic has  $\dot{u} = \dot{\eta} = 0$ ; thus its pair vector  $\xi$  is spatial-only and  $PI_X P$  discards all  $u, \eta$  components. Zero dissipation is consequently equivalent to  $A_X h(v) - A_X h(v')$  being parallel to  $V_X(v) - V_X(v')$  for all local velocity pairs. Positivity of the pair weight first gives the statement almost everywhere; smoothness extends it throughout the interior.

The radial component gives  $h_z$  independent of  $(u, \mu)$ . The map  $(u, \mu) \mapsto (u, (mu^2 + \mu C/R)/(qC))$  is invertible onto an open two-dimensional set. On such a set, a map whose every chord is parallel to the corresponding argument chord has the form  $F(v) = \alpha v + c$  with scalar  $\alpha$  constant on that velocity set. To prove this elementary lemma, fix one point and compare its chords with two noncollinear chords: parallelism of the third chord forces their scalar factors to agree. An open set supplies a noncollinear comparison for each additional point. Applied to the angular and vertical components of (55), this gives

$$\frac{h_u}{m} = \alpha(\mathbf{X})u + c_\theta(\mathbf{X}), \quad \frac{-Rh_R + uh_u}{qC} = \alpha(\mathbf{X})\frac{mu^2 + \mu C/R}{qC} + c_z(\mathbf{X}). \quad (56)$$

Integrate the first identity in  $u$  to obtain  $h = m\alpha u^2/2 + mc_\theta u + d(\mathbf{X}, \mu)$ . Comparing the  $u^2$ ,  $u$  and constant terms in the second yields

$$\alpha_R = 0, \quad (c_\theta)_R = c_\theta/R, \quad d_R = -\alpha\mu C/R^2 - qC c_z/R. \quad (57)$$

Hence  $c_\theta = Rb(\theta, z)$  and  $h = \alpha(\theta, z)E + b(\theta, z)mRu + g(\theta, z, \mu) + \phi(R, \theta, z)$ . The already obtained velocity independence of  $h_z$  forces  $\alpha_z = b_z = 0$  and makes  $g_z$  independent of  $\mu$ ; its remaining velocity-independent part can be absorbed into  $\phi$ . This proves necessity of (53). The product velocity domain ensures the integration constant is a globally single-valued function of  $\mu$ ; connected support alone would not exclude branchwise functions on disconnected  $u$ -fibers.

For sufficiency, a spatial-only  $\phi$  has velocity-independent spatial action, a  $g(\theta, \mu)$  has zero spatial action, an  $a(\theta)E$  has pair spatial action  $a\xi$ , and a  $b(\theta)mRu$  has spatial action  $b(\theta)R\mathbf{e}_\theta$  independent of velocity. In the last identity,  $\mathbf{b} \times (m\mathbf{u}\mathbf{e}_R)/(qB)$  cancels the curvature part from  $\partial_u(mRu) = mR$ . Their projected pair differences therefore all vanish. This also proves that all spatial populations are conserved by local collisions, including the nonlinear surrogate. Since the local energy-flow map is injective, distinct toroidal velocity states never have  $\xi = 0$ ; the coincident diagonal has vanishing test-function difference.

**Proof of the common stationary statement.** The ideal derivative is  $\mathcal{T}h = (u/R)h_\theta + [(mu^2 + \mu C/R)/(qC)]h_z$ . Substitute (53) and regard it as a polynomial in  $u$ . The cubic coefficient forces  $a' = 0$ ; the constant coefficient  $\mu\phi_z/(qR)$  forces  $\phi_z = 0$ , since the moment interval has nonzero length. The quadratic coefficient then forces  $b' = 0$ . The remaining linear coefficient is  $\phi_\theta + g_\theta = 0$ ; angle-only redundant terms can be absorbed between  $\phi$  and  $g$ , giving (54). Periodic  $z$  supplies the closed boundary for the norm argument, but is not needed for this pointwise algebraic intersection. A proposed nonperiodic compensator  $\phi = -qCb'(\theta)z$  cancels a quadratic term but leaves  $-\mu Cb'/R$ , and is not stationary on open moment support.

Thus combined dynamics retain the full radial population, the full moment marginal, energy and toroidal canonical momentum  $mRu$ . A general joint  $(R, \mu)$  marginal is not implied. Arbitrary radial profiles cannot approach a unique state determined only by energy and  $G(\mu)$ . The same argument classifies positive nonlinear collision equilibria by taking  $h = \ln f$  when the entropy form is finite; their common ideal stationary family is  $\ln f = \phi(R) + g(\mu) + aE + bmRu$ . Open boundaries with imposed flux, velocity support of dimension one, vanishing pair weights, disconnected velocity fibers, a nonlocal spatial kernel, or an electrostatic field that changes the energy-flow components are outside this proposition.

On the  $3^5 = 243$ -node test, 27 spatial indicators, nine angular-moment indicators, three angular-energy functions and three angular-momentum functions have rank 39 after three duplicated angle-only functions are counted once. Their production actions vanish to roundoff and account for all 39 observed near-null eigenvalues of the unmodified matrix. The continuum proof above is independent of that finite-grid rank check; eigenvalues are not altered. A separate continuous-bracket check at  $m = 1.7$ ,  $q = -0.8$  and three off-axis points verifies random classified functions and the cylindrical action formula.

Primary-source searches of the authors with toroidal collision nullspace, radial populations and moment constraints, and local review of SM25 and BS25, did not locate this specific local-kernel classification. Local density and symmetry-generated invariants are established ideas; publication priority for this closure-specific classification remains unresolved. No claim about physical Coulomb transport, general magnetic geometry, or a measured lifetime follows from it.

## Nonzero-field bounds for the executed domains

For the collision-only Cartesian box  $0.8 \leq x \leq 1.2$ ,  $|y| \leq 0.2$ ,  $0.1 \leq z \leq 0.5$ , the mirror with  $B_c = 1$ ,  $a = 0.15$  has  $B \geq B_z \geq 1 + 0.15[0.1^2 - (1.2^2 + 0.2^2)/2] = 0.8905$ . For the unit dipole,  $B = r^{-3}\sqrt{1 + 3z^2/r^2} \geq (1.2^2 + 0.2^2 + 0.5^2)^{-3/2} > 0.439$ . The symmetry-breaking addition with  $\varepsilon = 0.03$ ,  $B_c = L = 1$  has norm at most  $0.03\sqrt{1.2^2 + 0.2^2} < 0.037$ , so the perturbed dipole remains nonzero throughout that box, with  $B > 0.402$ . These conservative analytic bounds certify the geometry over the domain, beyond sampled minima. They do not certify a guiding-center ordering or a confined Hamiltonian boundary. For the toroidal annulus  $1 \leq R \leq 1.5$ ,  $C = 1$ ,  $B \geq 2/3$  exactly.

## 10 Vacuum energy-flow zero sets and spatial population constraints

The results in this section concern the prescribed local scalar  $DPI_X P$  closure, a time-independent nonzero vacuum field, constant electrostatic potential,  $m > 0$ ,  $q \neq 0$ , and the common Cartesian  $\eta$  chart. The positive scalar  $D$  is a symmetric local pair weight and may depend on the local velocity pair. Statements about the nullspace require it and the equilibrium density to be strictly positive almost everywhere. Global conservation statements additionally require the closed weak collision boundaries used in the entropy identity. These are continuous statements; a modified discrete energy gradient is a separate approximation.

### Injectivity away from critical magnetic strength

Set  $\mathbf{g} = \nabla B$ ,  $\mathbf{t} = \mathbf{b} \times \mathbf{g}$  and  $s = \mathbf{b} \cdot \mathbf{g}$ . The vacuum identity  $\nabla \times \mathbf{b} = \mathbf{t}/B$  follows by differentiating  $\mathbf{b} = \mathbf{B}/B$  and using  $\nabla \times \mathbf{B} = 0$ . For  $E = mu^2/2 + \mu B$ , the energy flow in the common chart is exactly

$$V = T\mathcal{J}_\mu \nabla E = \left( u\mathbf{b} + \frac{mu^2 + \mu B}{qB^2} \mathbf{t}, -\frac{\mu s}{m}, \mu s \right). \quad (58)$$

Here  $mu^2$  denotes mass times  $u^2$ , whereas  $\mu$  is the magnetic moment. At fixed spatial position,  $\mathbf{b} \cdot V_X = u$ . Equality of two flow vectors therefore fixes  $u$ . If  $\mathbf{t} \neq 0$ , the remaining spatial component fixes  $\mu$ . If  $\mathbf{t} = 0$  but  $s \neq 0$ , the parallel-velocity component  $V_u = -\mu s/m$  fixes  $\mu$ . Thus the map  $(u, \mu) \mapsto V$  is injective whenever  $\nabla B \neq 0$ . Its derivative has rank two: the  $u$  column has  $\mathbf{b} \cdot \partial_u V_X = 1$ , while the  $\mu$  column has zero such projection and is nonzero. Conversely, at  $\nabla B = 0$  the map is  $V = (u\mathbf{b}, 0, 0)$ , so distinct equal- $u$  moment states have equal flows. This classifies the fixed-position zero set without regularizing  $|\xi|^2$ . The fixed-position qualification is essential. In the unit dipole with  $m = q = 1$ , the distinct equatorial states  $(\mathbf{X}, u, \mu) = ((1, 0, 0), 0.2, 0.5)$  and  $((2, 0, 0), 0.2, 0.76)$  both have common-chart energy flow  $(0, 1.62, -0.2, 0, 0)$ , even though  $\nabla B$  is nonzero at both positions. At the equator  $B = C/R^3$ ,  $\mathbf{b} = -\mathbf{e}_z$ , and the azimuthal flow is  $3[mu^2 R^2/C + \mu/R]/q$ . This distinct remote pair also has a nonremovable directional weak limit for the spatial scalar projector. For  $h = u$ ,  $(Au)_X = -\mathbf{e}_z + 3uR^2\mathbf{e}_\theta$ , so the pair action difference is  $-1.8\mathbf{e}_\theta$ . Perturbing only  $\mu$  at the first state makes  $\xi$  parallel to  $\mathbf{e}_\theta$  and the weak factor vanish. Perturbing only  $u$  by  $\delta$  instead gives

$$|I_X P \Delta Au|^2 = \frac{(1.8 - 3\delta)^2}{1 + (1.2 + 3\delta)^2} \rightarrow \frac{3.24}{2.44} = 1.327868852 \dots \quad (59)$$

The actual common-chart action and projection are tested along three nonsingular values of each perturbation; the singular pair is not assigned a projector. Thus local injectivity supplies no general equal-flow exclusion or continuous zero-vector prescription for a nonlocal scalar extension. This observation alone does not analyze the full interaction tensor in the nonlocal source Eq. (119).

The point dipole has no critical strength point away from its source: homogeneity gives  $\mathbf{X} \cdot \nabla B = -3B$ . For the mirror, writing  $B_z = C + a(z^2 - R^2/2)$  gives

$$\partial_R(B^2) = 2aR(-C + aR^2/2), \quad \partial_z(B^2) = 4az(C + az^2). \quad (60)$$

For  $C > 0$  and  $a \neq 0$ , simultaneous zeros in the nonzero-field region occur only at the origin. The executed positive- $z$  box excludes it. For a dipole plus the degree-one harmonic perturbation, write  $\mathbf{B} = \mathbf{B}_d + \mathbf{B}_p$ . Euler homogeneity yields

$$\frac{1}{2} \mathbf{X} \cdot \nabla(B^2) = \mathbf{B} \cdot (-3\mathbf{B}_d + \mathbf{B}_p) \leq -3|\mathbf{B}_d|^2 + 2|\mathbf{B}_d||\mathbf{B}_p| + |\mathbf{B}_p|^2 < 0 \quad (61)$$

whenever  $|\mathbf{B}_p| < |\mathbf{B}_d|$ . The bounds above establish this on the executed perturbed-dipole box. Hence all three nonuniform Cartesian test domains have injective continuous local energy-flow maps. This does not by itself certify injectivity of a finite-grid energy-flow approximation.

For a smooth test function,  $\Delta Ah = O(|\Delta(u, \mu)|)$  near a coincident pair. Since  $PI_X P$  is bounded, its weak quadratic factor is  $O(|\Delta(u, \mu)|^2)$  and extends uniquely by zero on the diagonal, even though the projector itself has a directional limit. In contrast, a distinct equal-flow pair at a critical spatial point need not have a continuous weak extension.

## A nonremovable critical-point weak limit

Consider the mirror near its nonzero-field origin, equal  $u \neq 0$ ,  $\mu \neq \mu'$ , and  $h = \mu x$ . Approach along the  $x$  axis with  $z = 0$ . Then  $\mathbf{b} = \mathbf{e}_z$ ,  $\mathbf{t}$  is parallel to  $\mathbf{e}_y$ , and both  $\xi$  and  $\Delta Ah$  are spatial and parallel. Therefore  $|I_X P \Delta Ah|^2 = 0$ . Approach instead along the  $z$  axis. Now  $\mathbf{t} = 0$  and  $\xi$  has only  $u$  and  $\eta$  components, while  $\Delta Ah = (\mu - \mu')\mathbf{e}_y/(qB)$ . Consequently,

$$\lim_{x \rightarrow 0, z=0} |I_X P \Delta Ah|^2 = 0, \quad \lim_{z \rightarrow 0, x=0} |I_X P \Delta Ah|^2 = \frac{(\mu - \mu')^2}{q^2 C^2} > 0. \quad (62)$$

Varying  $u - u'$  at the exact origin gives the latter uniform-field prescription. Thus a simultaneous small-gradient and equal- $u$  limit is not equivalent to first imposing the uniform model. There is no universal continuous extension at these distinct-state critical pairs. An arbitrary bounded assignment on a measure-zero critical spatial set does not change an absolutely continuous weak integral, but it cannot be described as its unique continuous limit. Positive quadrature nodes on such a set need a separately declared prescription. No denominator softening is justified by this argument.

**Near-uniform mirror fields and tensor diagonal weights.** Example 18 tests a finite-quadrature limitation using exact continuous local actions, independently checked against the common-chart implementation. Its observable is the frozen positive pair Gram  $Q_h$  for  $h = \mu x$ , evaluated at  $f = e^{-E-0.2\mu}$ ; it is not this equilibrium's entropy production, which is zero. At  $X = (1, 0, 0.3)$ ,  $m = q = C = D = 1$ , the positive-amplitude mirror field has  $\Delta Ah \rightarrow (\mu - \mu')\mathbf{e}_y$ . For  $u \neq u'$ ,  $\Delta V \rightarrow (u - u')\mathbf{e}_z$  and the projector approaches the literal uniform prescription. For  $u = u'$  and  $\mu \neq \mu'$ , however,

$$\Delta V = a(\mu - \mu')L(u) + O(a^2), \quad L(u) = (0, -1, 0, -0.6, 0.6u).$$

Consequently the limiting spatial squared projected action is  $(\mu - \mu')^2 r(u)$ , with  $r(u) = [0.36(1 + u^2)/(1 + 0.36(1 + u^2))]^2 < 1$ , whereas the uniform prescription gives  $(\mu - \mu')^2$ . For positive tensor quadrature weights  $w_i$  in  $u$  and  $v_j$  in  $\mu$ , put  $M_k = \sum_j v_j e^{-1.2\mu_j} \mu_j^k$  and  $H = M_0 M_2 - M_1^2 > 0$ . The unordered pair sum gives exactly

$$Q_n(0) - \lim_{a \rightarrow 0^+} Q_n(a) = H \sum_i w_i^2 e^{-u_i^2} [1 - r(u_i)] \leq H \max_i w_i \sum_i w_i e^{-u_i^2}.$$

Thus a fixed quadrature retains a positive discrepancy. Positive Gauss refinement has  $\sum_i w_i = 2U_{\max}$  and  $\max_i w_i \rightarrow 0$ , so the displayed bound vanishes when the moment rule is held fixed. On the fixed compact velocity domain the projector and analytic test actions are bounded. Dominated convergence therefore recovers the same uniform continuum form because  $u = u'$  has zero product measure. This is a finite-quadrature discontinuity, not a physical discontinuity or a claim that the two infinite sequential limits differ. The example holds the moment quadrature and velocity tails fixed, isolates the parallel-order effect, and reports the direct positive-amplitude checks and proved bound. It does not certify the modified discrete-energy projector or a full spatial evolution scheme uniformly as  $a \rightarrow 0$ , and makes no priority claim for this numerical limitation.

## Spatial invariants of local collisions

For a spatial observable  $\phi(\mathbf{X})$  independent of  $u, \mu$ , its common-chart action is

$$A\phi = \left( \frac{\mathbf{b} \times \nabla \phi}{qB}, -\frac{\mathbf{b} \cdot \nabla \phi}{m} - \frac{u \mathbf{t} \cdot \nabla \phi}{qB^2}, -\frac{\mu \mathbf{t} \cdot \nabla \phi}{qB} \right). \quad (63)$$

If  $\mathbf{t} \cdot \nabla \phi = 0$ , this entire vector is independent of the velocity pair at the same spatial position, so  $\Delta A\phi = 0$ . Every such observable is conserved by the linear and nonlinear local collision surrogates, regardless of the energy projector.

There is also a precise local necessity statement for spatial observables. At a point with  $s \neq 0$  and  $\mathbf{t} \neq 0$ , let  $\kappa = \mathbf{t} \cdot \nabla \phi$ . Choose equal nonzero  $u$  and unequal  $\mu, \mu'$  in a product velocity domain of nonzero widths. Their action difference has only an  $\eta$  component, while  $\xi_X = (\mu - \mu')\mathbf{t}/(qB)$  and  $\xi_\eta = us(\mu - \mu')$ . If  $\kappa \neq 0$ , then  $\xi \cdot \Delta A\phi = -us\kappa(\mu - \mu')^2/(qB) \neq 0$  and the spatial block of  $P\Delta A\phi$  is nonzero. Continuity gives a positive weak contribution on a neighborhood of this pair. Thus a spatial observable is locally null precisely when  $\kappa = 0$  at such a point. If  $s = 0$ , the energy-flow difference is spatial, and the velocity-dependent part of  $\Delta A\phi$  is nonspatial; every spatial observable is locally null. The case  $\mathbf{t} = 0$  already has  $\kappa = 0$ . These statements classify spatial test functions, not the full nonuniform nullspace.

Locally a vacuum field has  $\mathbf{B} = \nabla \chi$ . Since  $\mathbf{t}$  is orthogonal to both  $\nabla B$  and  $\nabla \chi$ , every spatial  $F(B, \chi)$  is a collision invariant. When  $\chi$  is globally single valued, closed collisions therefore preserve the joint spatial pushforward population in  $(B, \chi)$ . The executed scalar potentials are

$$\chi_m = Cz + a(z^3/3 - zR^2/2), \quad (64)$$

$$\chi_d = -Cz/r^3, \quad \chi_{na} = \chi_d + \varepsilon Cxy/L. \quad (65)$$

Ideal flow generally changes these populations. They are not interior Casimirs of the Poisson tensor and are not claimed to be invariants of the full combined dynamics.

## Poloidal flux memory and inaccessible dipole candidates

In an axisymmetric purely poloidal field,  $\mathbf{t}$  is azimuthal. Therefore all  $\phi(R, z)$  satisfy the preceding sufficient condition: local collisions preserve every meridional spatial population. With  $\mathbf{B} = \nabla \psi \times \mathbf{e}_\theta/R$ , the poloidal fluxes are

$$\psi_d = \frac{CR^2}{(R^2 + z^2)^{3/2}}, \quad \psi_m = \frac{CR^2}{2} + \frac{aR^2z^2}{2} - \frac{aR^4}{8}. \quad (66)$$

Both  $\mathbf{b}$  and  $\mathbf{t}$  are tangent to these flux surfaces. For constant potential, ideal flow conserves  $\psi$ ; the same remains true for an axisymmetric electrostatic potential. Under closed, flux-compatible ideal boundaries and closed weak collision boundaries,

$$\frac{d}{dt} \int F(\psi) f \, d\Gamma = 0 \quad \text{for every admissible smooth } F. \quad (67)$$

This is conservation of the full flux pushforward marginal, which obstructs accessibility of an energy/ $G(\mu)$ -only equilibrium from arbitrary initial data. It does not generally preserve a joint  $(\psi, \mu)$  population: for example  $h = \mu\psi$  has a nonzero projected pair action off the equatorial plane. The test suite checks that counterexample explicitly. In particular, under the same closed combined conditions, the stationary family includes

$$f = \exp[-\beta E + g(\mu) + F(\psi)] \quad (68)$$

whenever it is positive, normalizable and in the weak form domain. The energy term is annihilated by  $P$ , the moment term by  $\mathcal{J}$ , and the flux term has zero local pair difference; all three terms are ideal constants. The earlier energy/moment candidate is therefore a proper subset of this stationary family, not a complete equilibrium classification for this field and closure. No completeness statement for (68) is asserted.

Canonical toroidal momentum conservation and the ideal flux constraint in the purely poloidal limit have established antecedents [9]. Here  $A\psi = (R\mathbf{e}_\theta/q, 0, 0)$  with the flux convention in (66); thus  $q\psi$  is the lowest-order canonical rotation generator. The additional conclusion is that exact spatial localization makes *every*  $F(\psi)$  collisionally invariant. This proof uses equality of the paired positions essentially. It does not apply to the generally spatially nonlocal source Eq. (119), to its full interaction tensor, or to finite-Larmor collision displacements. The source introduces the spatial scalar/delta approximation separately in Eqs. (171)–(172). Conventional guiding-center collision operators retain mechanisms absent from this closure [4, 7]. Flux- or canonical-momentum-dependent equilibria are not new objects: Di Troia [10] treats equilibria depending on canonical toroidal momentum, energy and moment, and modifies the Landau projection using the selected equilibrium score difference in Eqs. (44) and (46) of the checked arXiv version. That target-dependent construction differs from a nullspace constraint forced by the present fixed local pair geometry. Sato, Yoshida and Kawazura [11] construct inward diffusion from fluctuating electric fields on a moment-preserving leaf; their Eq. (14d) contains a nonzero Wiener increment in  $\psi$ , and Eq. (15) contains flux diffusion. Such forcing changes the flux population and lies outside the closed axisymmetric ideal-plus-local-collision assumptions here. These antecedents preclude a broad novelty claim for constrained equilibria or dipole transport. The present negative result concerns the specified closure’s inaccessible states, not a general prohibition of inward diffusion. Focused primary-literature searches found no matching statement for this precise surrogate; no publication-priority claim follows. Neither the invariant theorem nor the critical-point counterexample provides a physical rate, a quantitative decay bound or a dipole lifetime.

## Symmetry breaking removes spatial ideal constants

For the unit dipole plus  $a\nabla(xy)$ ,  $0 < |a| \leq 0.03$ , on the connected Cartesian box stated above, there is no nonconstant  $C^2$  spatial function  $\phi(\mathbf{X})$  that is an ideal constant for every velocity state in an open product domain. Consequently there is no nonconstant spatial function in the common ideal/collision nullspace. This concerns spatial observables only; it does not classify functions depending on  $u, \mu$ , establish coercivity, or imply a relaxation rate.

Indeed, comparing the  $\mu$  and  $u$  coefficients of  $V_X \cdot \nabla \phi = 0$  in (58) gives  $\mathbf{b} \cdot \nabla \phi = \mathbf{t} \cdot \nabla \phi = 0$ . Set  $S = |\mathbf{B}|^2$ ,  $p = \mathbf{B}$ , and  $v = \mathbf{B} \times \nabla S = 2B^2\mathbf{t}$ . Both vector fields annihilate  $\phi$ , so their Lie bracket  $[p, v] = (\nabla v)p - (\nabla p)v$  also annihilates it. At the interior unit-radius point  $(15/17, 0, 8/17)$ , exact rational differentiation gives

$$\det[p, v, [p, v]] = \frac{900a(225a^2 + 481)(6709a^4 - 313674a^2 + 2686977)}{6975757441}. \quad (69)$$

The quartic factor is at least  $2686977 - 313674(0.03)^2 > 2686694$  on the stated amplitude interval. Thus the determinant is nonzero. The field is real analytic and bounded away from zero throughout this box, and the analytic determinant is not identically zero. Its nonzero set is therefore dense. On that set the three annihilating vector fields span physical space, so  $\nabla \phi = 0$ ; continuity and connectedness make  $\phi$  constant throughout the box. For  $a = 0.03$ , the exact value is  $1396271742281072028423/279030297640000000000 = 5.004014811619193\dots$ ; direct field automatic differentiation gives 5.004014811619226.

The certificate was derived independently by exact rational field derivatives and by symbolic differentiation; regression tests check the actual field implementation at both perturbation signs. This two-vector Lie-rank argument is distinct from the one-field Pfaff/Frobenius condition  $\mathbf{B} \cdot \nabla \times \mathbf{B} = 0$  discussed in SM24 [2]. The latter holds for all these vacuum fields. The result removes the axisymmetric spatial flux constraint, not all phase-space constraints or all possible obstructions to mixing. Focused primary-source searches found no matching statement for this fixed analytic example, but no publication-priority claim is made.

## 11 Linear weak discretization and SOLVAX

Represent  $h = \sum_i a_i \chi_i$ . Define the mass and collision matrices by

$$M_{ij} = \int f_0 \chi_i \chi_j d\Gamma, \quad (70)$$

$$K_{ij} = \frac{1}{2} \iint f_0 f'_0 \Delta \chi_i^T \Pi_0 \Delta \chi_j d\Gamma d\Gamma'. \quad (71)$$

Positive quadrature gives a Gram form  $K = B_p^T W_p B_p \succeq 0$ . Implement its action without assembling the full pair tensor. A dense assembly on a tiny grid must remain an independent reference.

The executable nodal construction uses tensor differentiation followed by  $A_i \phi = T_i \mathcal{J}_{\mu,i} (\nabla_h \phi)_i$  in a common Cartesian  $\eta$  chart. For every unordered pair  $(i, j)$ , set  $B_{ij} \phi = A_i \phi - A_j \phi$  and construct  $\xi_{ij} = B_{ij} e$  from the same discrete energy vector  $e$  and derivative. The nonuniform projector uses this discrete  $\xi$ , rather than inserting an incompatible continuum energy gradient. This is an explicitly stated discrete energy projector; its consistency must be checked under refinement. The scalar pair weight is  $w_X (B w_v)_i (B w_v)_j$  locally and  $w_i w_j W_{ij}$  for a spatially nonlocal model. In a toroidal chart the spatial weight contains  $R dR d\theta dz$  while each local velocity factor contains  $B d\mu d\mu$ , even though the one-particle product  $BR = C$  simplifies.

The ordered weak form in (71) counts  $(i, j)$  and  $(j, i)$  equally: exchange reverses both pair differences, leaving their quadratic product unchanged. Therefore its  $1/2$  exactly cancels the two ordered contributions, and each unordered pair is inserted once with no  $1/2$ . Self-node pair differences are identically zero and need no stored contribution. Distinct nodes with equal energy-flow vectors are a separate issue: the uniform field uses its derived continuous limit, while an exactly zero nonuniform pair vector is rejected. No  $|\xi|^2 + \epsilon$  is used. Directional limits may depend on the approach. The continuous local zero set is classified by (58); the critical-point counterexample (62) prevents a universal continuous prescription. Successful finite-grid examples do not by themselves establish a valid extension for a modified discrete energy map.

Periodic derivatives are Fourier derivatives; the real even-grid Nyquist component is unresolved and given zero derivative. Nonperiodic derivatives support local three-point quadratic interpolation and, for the refined nonuniform boxes, the derivative of the full nodal Lagrange polynomial. The derivative and its transpose enter the weak collision assembly. The selected spatial quadrature and derivative method are recorded explicitly. This imposes the natural collision flux rather than wrapping nonperiodic field values. The preassembled reference path retains pair indices and directions with quadratic velocity-pair storage (and spatial pairs for nonlocal models). The compact local path instead stores energy flows and forms pair directions only inside bounded blocks, as described below.

If  $c$  represents a collision invariant and the derivative/kernel construction respects it, then  $Kc = 0$ . Include the discrete energy and all resolved magnetic-moment-bin constraints. Finite differences do not automatically obey the product rules needed for an arbitrary sampled  $\mu B$ . Use the same discrete

energy gradient in the degeneracy direction and in the weak operator, and prove consistency as the mesh is refined.

The collision-only system and norm balance are

$$M\dot{a} = -Ka, \quad \frac{d}{dt} \frac{1}{2} a^T Ma = -a^T Ka \leq 0. \quad (72)$$

With a compatible stationary Hamiltonian discretization,

$$M\dot{a} = (A - K)a, \quad A^T = -A. \quad (73)$$

Also test  $Ac = 0$  for the required linear invariants; skew symmetry alone does not imply those individual identities.

In the nonuniform combined example, exact Fourier half-translations in the periodic toroidal angle and  $z$  surround one backward-Euler collision step. The translation uses  $\dot{\theta} = u/R$  and  $\dot{z} = (mu^2 + \mu C/R)/(qC)$  at each fixed  $(R, u, \mu)$ . The measure  $BR = C$ , energy and moment marginal are invariant under this ideal flow, and radial walls are tangent. This symmetric composition remains first order because the collision substep is backward Euler. It is a linear, fixed-equilibrium combined evolution; the nonlinear examples are collision only.

A backward-Euler collision step solves

$$(M + \Delta t K)a^{n+1} = Ma^n. \quad (74)$$

For positive-definite  $M$ , this is an appropriate PCG problem. SOLVAX's `pcg_linear_solve` supplies a converged-system differentiation wrapper [12]. Do not pass  $M^{-1}K$  to Euclidean PCG without mass scaling or a proof of symmetry. A coupled implicit step with  $A$  is generally nonsymmetric and calls for GMRES or another suitable method.

For a linear invariant  $c$ , multiply (74) by  $c^T$  to check its conservation. Backward Euler also satisfies

$$Q^{n+1} - Q^n = -\Delta t (a^{n+1})^T K a^{n+1} - \frac{1}{2} \|a^{n+1} - a^n\|_M^2 \leq 0. \quad (75)$$

This does not prove positivity of a reconstructed finite-amplitude distribution. An implicit midpoint step can be second order, but is not automatically a good stiff-decay scheme or a nonlinear entropy-preserving integrator.

Local pair application costs  $O(N_X N_v^2)$  in a direct quadrature. A spatially nonlocal direct application can cost  $O(N_X^2 N_v^2)$ . Chunk pair contractions, reuse geometry and quadrature, and keep derivatives sparse. Block-tridiagonal solves may be useful in a line preconditioner; the full velocity coupling is not generally tridiagonal. The current tensor derivative matrices are stored densely, including zeros outside the nonperiodic stencils. Their contraction adds  $O(N \sum_\alpha n_\alpha)$  work, with  $N = \prod_\alpha n_\alpha$ ; stencil sparsity is a mathematical property, not an implemented sparse matmul claim.

## 12 A nonlinear discrete entropy construction

This finite-dimensional fixed-field construction is implemented. It is not a guarantee for arbitrary interpolation, self-consistent fields, or nonlinear Hamiltonian transport. It follows the weak-structure approach developed for Landau discretizations [13, 14, 15, 16].

Let  $f_i > 0$ ,  $w_i > 0$ , and  $n_i = w_i f_i$ . Construct a symmetric positive mobility  $K(f)$  from the discrete pair weak form, with

$$K\mathbf{1} = 0, \quad Ke = 0, \quad Kc_\mu = 0, \quad (76)$$

where  $e$  is the fixed-field energy vector and  $c_\mu$  spans the represented moment-marginal constraints. Define

$$S(n) = - \sum_i n_i \ln(n_i/w_i). \quad (77)$$

Then the collision ODE can be written

$$\dot{n} = K(f) \nabla_n S = -K(f) \ln f. \quad (78)$$

The constant term in  $\nabla_n S = -\ln f - \mathbf{1}$  vanishes through (76). Consequently  $\dot{S} \geq 0$  and  $d(e^T n)/dt = 0$ . This is an algebraic guarantee only for the actual constructed  $K$ .

For each component define a divided difference of  $s_i(n) = -n \ln(n/w_i)$ :

$$(\bar{\nabla} S)_i = \begin{cases} [s_i(n_i^+) - s_i(n_i^-)] / (n_i^+ - n_i^-), & n_i^+ \neq n_i^-, \\ -\ln(n_i^-/w_i) - 1, & n_i^+ = n_i^-. \end{cases} \quad (79)$$

Evaluate it stably near equal arguments using logarithmic divided differences, not cancellation-prone subtraction. A step

$$n^+ - n^- = \Delta t \bar{K} \bar{\nabla} S, \quad \bar{K} = \bar{K}^T \succeq 0 \quad (80)$$

with the null directions in (76) gives

$$S(n^+) - S(n^-) = \Delta t (\bar{\nabla} S)^T \bar{K} \bar{\nabla} S \geq 0. \quad (81)$$

The conservation laws follow by left multiplication with each invariant vector. Use a positive mean state to construct  $\bar{K}$  and a converged nonlinear solve. Parameterizing  $n_i^+ = w_i e^{g_i^+}$  enforces positive nodal populations, but does not prove existence of the solve or positivity between interpolation nodes. Check all required quadrature/reconstruction locations.

## Angular leaf components of the Cartesian box

The mirror and dipole collision-only boxes have an additional domain qualification. In the open rectangle  $0.8 < x < 1.2$ ,  $|y| < 0.2$ , a fixed circle with  $1.2 < R < \sqrt{1.48}$  meets the box in two disconnected angular arcs, with  $\arccos(1.2/R) < |\theta| < \arcsin(0.2/R)$ . For  $g \in C_c^\infty((1.2, \sqrt{1.48}))$ , the observable  $\phi = g(R) \text{sign}(y)$  is smooth inside the box: its active support is separated from  $y = 0$ . Since  $\mathbf{t}$  is azimuthal,  $\mathbf{t} \cdot \nabla \phi = 0$ . Closed local collisions therefore conserve these signed component populations as well as the single-valued meridional populations. This is a consequence of the natural closed boundary, not a failure of that boundary condition.

The corner region has unweighted geometric area fraction

$$\frac{2}{0.16} \int_{\sqrt{1.4}}^{1.2} [0.2 - \sqrt{1.44 - x^2}] dx = 0.01394734125 \dots \quad (82)$$

Its volume fraction is the same because the  $z$  interval is independent; its state-measure or particle fraction requires the actual  $B$  and density weights. Five-point Gauss nodes per axis have maximum  $R = 1.19506$  and miss this region; seven-point nodes reach  $R = 1.20487$  with two corner nodes in the  $xy$  grid. Thus independent bulk-observable refinements do not by themselves certify preservation of every localized leaf-component constraint or global continuum convergence. This analytical limitation is reported independently of the finite-time outcomes; it introduces no later acceptance criterion.

## Matrix-free Newton–GMRES and residual control

The compact nonuniform evolution solves the same discrete-gradient equation as the small dense reference. Let  $l = \ln f^+$ ,  $l^- = \ln f^-$ ,  $\bar{f} = (e^l + e^{l^-})/2$ , and write the unscaled residual as

$$F(l) = w(e^l - f^-) - \Delta t K(\bar{f}) \bar{\nabla} S(l^-, l). \quad (83)$$

Products with  $w$  are componentwise. Set  $a_i = \sqrt{w_i f_i^-} > 0$ ,  $r = F/a$ , and parameterize the Newton correction as  $\delta l = v/a$ . The exact matrix-free operator is

$$v \mapsto a^{-1} J_F(l)(v/a), \quad a^{-1} J_F(l)(v/a) = -r. \quad (84)$$

For the compact path, the pair tangent is prepared exactly. If  $\mathcal{A}h$  is the nodal common-chart action, a pair contribution to  $K(\bar{f})h$  is proportional to  $\bar{f}_i \bar{f}_j P I_X P \Delta \mathcal{A}h$ . Its variation is

$$\delta[K(\bar{f})h] = \mathcal{A}^T \sum_{ij} D w_X v_i v_j \bar{f}_i \bar{f}_j P I_X P [\Delta \mathcal{A} \delta h + (r_i + r_j) \Delta \mathcal{A} h], \quad r_i = \frac{f_i^+}{f_i^+ + f_i^-} \delta l_i. \quad (85)$$

The sum denotes the same paired flux accumulation as the mobility, including the local velocity Jacobians  $v_i$ . Here  $h = \bar{\nabla} S$  and its componentwise derivative gives  $\delta h = h'_i \delta l_i$ . The actions  $\mathcal{A}h$ , nodal mean populations and scalar  $h'_i$  are cached once per Newton iterate; each Krylov matvec applies one projected pair tangent. No frozen-mobility approximation is made. Independent forward Jacobian-vector products on the explicit-pair path differentiate all terms and check this prepared tangent. SOLVAX restarted GMRES solves this system; a finite-step Newton Jacobian need not be symmetric positive definite, so PCG is not substituted. Neither route retains a linearization tape containing every primal pair block.

The nonlinear relative stopping norm is  $\|F/a\|_2 / \sqrt{\sum_i w_i f_i^-}$ , fixed by the previous state. The dense reference retains its original population residual norm; it solves the same unscaled equation with its own explicitly recorded tolerance. A backtracking line search decreases the chosen residual norm and rejects nonfinite or exponential-underflow candidates. It clips no population and applies no invariant correction. Since the divided difference gives  $\Delta S = \bar{\nabla} S \cdot w(f^+ - f^-)$ , the exact entropy defect is

$$\Delta S - \Delta t \bar{\nabla} S^T K(\bar{f}) \bar{\nabla} S = \bar{\nabla} S \cdot F, \quad |\text{defect}| \leq \|a \bar{\nabla} S\|_2 \|F/a\|_2. \quad (86)$$

The implementation checks this bound with a floating-point allowance. Likewise, if  $Kc = 0$ , the invariant error is  $c \cdot F$  and is bounded by  $\|ac\|_2 \|F/a\|_2$ . These are residual-dependent bounds, not exact conservation assertions for an unconverged root.

For the compact local Cartesian grid the stored coefficients are  $N \times 5 \times 4$ , and energy flows, populations, moment indices and local quadrature have  $O(N)$  entries. Dense axis derivative matrices add  $O(\sum_\alpha n_\alpha^2)$  storage. No local pair list or dense Newton Jacobian is stored. A target-row contraction processes all  $N_v$  partners; with pair budget  $C_p$ , temporary pair workspace is  $O(\max(C_p, N_v))$  fixed-dimensional vectors. Restarted GMRES additionally stores  $O(Nr)$  basis entries for restart length  $r$ . Thus the precise storage bound is  $O(N + \sum_\alpha n_\alpha^2 + Nr + \max(C_p, N_v))$ , apart from compiler/runtime storage. Measured grid bytes and process timings are reported independently. This does not supply an accelerated pair asymptotic: the direct local contraction still evaluates  $O(N_X N_v^2)$  interactions per matvec.

For the executed nonlinear uniform and linear toroidal examples, continuous spatial Fourier positivity is checked with a sufficient certificate. If  $f(x) = \sum_k \hat{f}_k e^{ikx}$ , then

$$f(x) \geq \hat{f}_0 - \sum_{k \neq 0} |\hat{f}_k|. \quad (87)$$

The same triangle bound holds for a double Fourier sum in  $(\theta, z)$ . The example rejects a nonpositive bound, and additionally oversamples the reconstruction. Positive piecewise-linear interpolation in the remaining coordinates then preserves positivity throughout that declared reconstruction. This certificate is conditional on the computed coefficients and is not a positivity theorem for arbitrary Fourier states or higher-order velocity interpolation. The collision derivative still uses its explicitly stated local quadratic differentiation, independently of the positivity-preserving display reconstruction. The Cartesian collision-only boxes use positive nodal values with piecewise-linear reconstruction.

For self-consistent fields the energy is nonlinear in  $n$ . Its discrete gradient must enter the collision degeneracy consistently; the fixed vector  $e$  above is not enough. Likewise, nonlinear Hamiltonian transport requires its own conservative and positivity-compatible treatment. A linear skew matrix alone is not a proof of a nonlinear H-theorem.

## Lagged mobility and an implicit entropy-variable step

For a finite positive-quadrature grid, let  $n_i = w_i f_i$  and freeze the constructed mobility at the old density:  $K_- = K(f_-)$ . Assume finite  $w_i, n_{-,i} > 0$ ,  $K_- = K_-^T \succeq 0$ ,  $K_- \mathbf{1} = 0$ , and  $\Delta t \geq 0$ . The implemented lagged-mobility step solves

$$n_+ - n_- = -\Delta t K_- g_+, \quad g_{+,i} = \ln(f_{+,i}/f_{\text{ref}}), \quad n_{+,i} = w_i f_{\text{ref}} e^{g_{+,i}}, \quad (88)$$

where  $f_{\text{ref}}$  is a fixed positive density scale. Computation uses the previously specified normalized units with  $f_{\text{ref}} = 1$ . A constant change of this scale drops out through  $K_- \mathbf{1} = 0$ . This changes the time discretization of the same nodal collision ODE, not the kernel, its boundaries or its physical calibration.

The residual  $R = w f_{\text{ref}} e^g - n_- + \Delta t K_- g$  is the gradient of

$$\Psi(g) = \sum_i w_i f_{\text{ref}} e^{g_i} - n_-^T g + \frac{\Delta t}{2} g^T K_- g. \quad (89)$$

Each exponential minus positive linear term tends to infinity at either end of the real line and has a finite lower bound. The finite sum is coercive; the nonnegative quadratic preserves this property. Moreover  $\nabla^2 \Psi = \text{diag}(n_+) + \Delta t K_- \succ 0$  at every finite  $g$ . There is therefore a unique finite minimizer and a unique positive root for every nonnegative timestep. This is a finite-dimensional existence statement, not a uniform conditioning or finite-precision convergence bound.

For  $\Delta t > 0$ , the equivalent primal problem is

$$\min_{n \geq 0, n - n_- \in \text{Range } K_-} \left\{ \sum_i n_i [\ln(n_i/(w_i f_{\text{ref}})) - 1] + \frac{(n - n_-)^T K_-^\dagger (n - n_-)}{2\Delta t} \right\}. \quad (90)$$

The range constraint retains every actual Gram null constraint. Since  $\mathbf{1}$  is null, the admissible populations have fixed total mass and form a compact set. Strict convexity gives a unique minimizer. A boundary minimizer is impossible: the feasible direction toward the strictly positive  $n_-$  has entropy derivative tending to  $-\infty$  at each zero component, while the quadratic derivative remains finite. At the interior minimizer, project the Euler equation onto  $\text{Range } K_-$  and multiply by  $K_-$  to recover (88). This is a fixed finite-dimensional Onsager metric; it is not an exact Wasserstein-distance or JKO step.

At an exact root,  $c^T(n_+ - n_-) = 0$  for every  $K_- c = 0$ . With a residual,  $c^T(n_+ - n_-) = c^T R$ . Positive frozen pair weights retain the same actual discrete Gram nullspace, including number, discrete energy and resolved moment-bin populations. They do not restore continuum constraints lifted by the spatial approximation. A represented stationary density remains stationary.

For  $S(n) = -\sum_i n_i \ln(n_i/(w_i f_{\text{ref}}))$ , define

$$\mathcal{D}(n_- \| n_+) = \sum_i [n_{-,i} \ln(n_{-,i}/n_{+,i}) - n_{-,i} + n_{+,i}] \geq 0. \quad (91)$$

Direct subtraction gives the exact budget, including a nonzero root residual,

$$S_+ - S_- = \Delta t g_+^T K_- g_+ + \mathcal{D}(n_- \| n_+) - (g_+ + \mathbf{1})^T R. \quad (92)$$

With  $a_i = \sqrt{n_{-,i}}$ , the residual entropy-defect bound is  $|(g_+ + \mathbf{1})^T R| \leq \|a(g_+ + \mathbf{1})\| \|R/a\|$ ; invariant defects have the analogous bound  $|c^T R| \leq \|ac\| \|R/a\|$ . The first positive term is discrete mobility dissipation approximating the continuous nodal model, without a calibrated physical rate. The KL term is a separate contribution from time discretization. For a smooth positive solution it is  $O(\Delta t^2)$  per step. Positivity and monotone entropy do not establish temporal accuracy or continuum convergence.

The implicit-function expansion at  $\Delta t = 0$  uses the invertible derivative  $\text{diag}(n_-)$  and gives  $n_+ - n_- = -\Delta t K(f_-) \ln(f_-/f_{\text{ref}}) + O(\Delta t^2)$ . Thus the step is first-order consistent with the same fixed spatial ODE. This local statement does not supply a mesh-uniform error bound or a convergence theorem near vanishing populations.

The positive Hessian permits PCG after entropy scaling, unlike the generally nonsymmetric finite-step discrete-gradient Jacobian. Candidate Newton directions must pass an objective-descent check, a true Hessian-residual check and objective backtracking. For a log increment  $s$ , use the stable exact objective change

$$\Psi(g + s) - \Psi(g) = \sum_i n_i (e^{s_i} - 1 - s_i) + R(g)^T s + \frac{\Delta t}{2} s^T K_- s. \quad (93)$$

The small-argument exponential remainder is evaluated by its series. Comparing independently evaluated objective values can erase the tiny Newton decrease near the root. No population clipping, invariant correction, nullspace removal or denominator softening is part of this step.

The method is described as lagged-mobility backward Euler in entropy variables, without a new-integrator claim. Shen and Xu [17] provide a close antecedent: their first-order scheme treats old density in the mobility and new logarithmic density implicitly, and their fully discrete positive-quadrature proof uses a PSD stiffness pseudoinverse and a constrained strictly convex entropy functional (Theorem 3.2, Eqs. 3.12–3.15). The present collision Gram matrix can have additional null directions, so the finite-dimensional argument above is stated independently; their PDE convergence results are not imported. General variational entropy-flow and metriplectic antecedents remain relevant context, but do not identify this frozen Gram metric with a Wasserstein distance.

**Implementation and independent checks.** The compiler accepts the old density dynamically and prepares the positive channel-lumped Jacobi approximation once per timestep. The approximation drops cross-derivative and pair cross terms only from the auxiliary inverse; the full Gram operator remains in the residual and Hessian. Original Hessian residuals are recomputed independently of recursive PCG stopping. The checked host step rejects nonfinite populations, failed linear corrections, non-descent directions, failed backtracking and inconsistent entropy budgets. A compiled operator is associated with its grid and solver configuration. For a derivative channel  $d$ , the positive auxiliary covariance is

$$q_{i,d} = D w_X v_i f_i \sum_{j \neq i} v_j f_j |I_X P_{ij} C_{i,d}|^2 \geq 0,$$

where  $v_i$  includes the velocity Jacobian and  $C_{i,d}$  is the corresponding common-chart action column. The diagonal approximation is  $\sum_d (D_d \circ D_d)^T q_d$ . The optional Cartesian  $x$ -line inverse instead retains

$L_x = D_x^T \text{diag}(q_x) D_x$  on each fixed  $(y, z, u, \mu)$  line. Its entropy-scaled block is

$$B_x = \text{diag}\left(\frac{n_+}{n_-} + \Delta t \frac{\sum_{d \neq x} (D_d \circ D_d)^T q_d}{n_-}\right) + \Delta t \text{diag}(n_-^{-1/2}) L_x \text{diag}(n_-^{-1/2}).$$

It is SPD for strictly positive finite populations: the first mass term is positive diagonal and the remaining terms are PSD. Batched Cholesky factors apply its inverse. This changes only PCG preconditioning; the full Hessian and all null constraints remain unchanged. Independent explicit-pair projectors check the covariance, and both inverse choices converge to the same separately assembled dense roots. No uniform conditioning bound follows. The campaign uses this line inverse and a 3000-iteration linear budget; failed residual checks remain fatal. API defaults retain the diagonal inverse and a 1000-iteration budget. For  $N$  state nodes and  $n_x$  nodes on a line, stored Cholesky factors use  $O(Nn_x)$  entries, their factorization costs  $O(Nn_x^2)$ , and each inverse application costs  $O(Nn_x)$ . This is additional solver storage beyond the earlier compact-grid arrays; the direct pair action still costs  $O(N_X N_v^2)$ .

Independent explicit-pair dense roots and a separately assembled pair ODE check tiny mirror, dipole and nonaxisymmetric grids. These method checks do not resolve the thermal velocity domain or establish a continuum limit.

### 13 Scattering physics and validity tests

The constrained model assumes negligible accumulated change of magnetic moment. That assumption is stronger than a small collision frequency compared with the gyrofrequency. At constant  $B$ , an energy change satisfies

$$\Delta E = \frac{m}{2} \Delta(u^2) + B \Delta\mu + q \Delta\Phi. \quad (94)$$

Small  $\Delta E$  does not separately bound  $\Delta\mu$ . In particular, the approximate identification in SM25 (85), printed p. 10, requires an extra bound on parallel-energy exchange: at constant  $B$ ,

$$B |\Delta\mu| \leq |\Delta E| + \frac{1}{2} m |\Delta(u^2)| + |q \Delta\Phi|. \quad (95)$$

Spatial locality and nearly constant  $B$  alone supply no bound on the second term. This audit concerns the inference from small energy change; it does not rule out a stronger regime in which every relevant exchange is ordered small. For independent unbiased increments, zero mean accumulation can coexist with

$$\text{Var}\left(\sum_{j=1}^N \delta\mu_j\right) = N \text{Var}(\delta\mu). \quad (96)$$

For independent increments, both  $N |\langle \delta\mu \rangle| / \mu_{\text{ref}} \ll 1$  and  $N \text{Var}(\delta\mu) / \mu_{\text{ref}}^2 \ll 1$  are needed for small accumulated change on a chosen nonzero scale. Compensating mean changes can still alter the full moment marginal through diffusion. These are audit points for the inference in SM25 (85) and the accumulation hypothesis (86), rather than replacements for its ordering. The source explicitly calls (84) and (86) working assumptions.

A controlled microscopic reference can integrate

$$m_s \dot{\mathbf{v}}_s = q_s \mathbf{v}_s \times \mathbf{B} - \nabla_{\mathbf{r}_s} V, \quad \dot{\mathbf{r}}_s = \mathbf{v}_s, \quad (97)$$

with a specified screened potential, for example  $V(r) = q_1 q_2 e^{-r/\lambda} / (4\pi\epsilon_0 r)$ . Screening is a declared model choice, not a result of this two-body calculation. Resolve close encounters or exclude them with

a stated grazing criterion. The generalized magnetized scattering theory in Ref. [18] is an important reference.

For uniform  $\mathbf{B}$ , compare with the unperturbed helical orbit at identical incoming conditions. Use

$$\mathbf{R}_s = \mathbf{r}_s + \frac{\mathbf{v}_s \times \mathbf{b}}{\Omega_s}, \quad \mu_s = \frac{m_s |\mathbf{v}_{\perp s}|^2}{2B}, \quad \Omega_s = q_s B / m_s. \quad (98)$$

Extract asymptotic incoming/outgoing guiding-center data at a common reference time. A raw endpoint displacement can contain an arbitrary flight-time contribution after the encounter changes parallel velocity.

Sample incident flux, impact parameters, gyrophases and relative velocities with the appropriate measure. Check timestep, start/end distance, angular sampling and screening/cutoff convergence independently. Recover a Rutherford limit before using strong-field results. Attractive close encounters outside the grazing regime are not automatically described by the reduced kinetic closure.

For independent encounters at a stated rate  $\nu_{\text{enc}}$ , diagnostic drift and diffusion coefficients are

$$A_\mu = \nu_{\text{enc}} \langle \delta\mu \rangle, \quad D_{\mu\mu} = \frac{\nu_{\text{enc}}}{2} \langle (\delta\mu)^2 \rangle. \quad (99)$$

Use both, together with cross-covariances and correlation tests. A timescale estimate must specify a nonzero population scale  $\mu_{\text{ref}}$ ; the quantities  $\mu_{\text{ref}}/|A_\mu|$  and  $\mu_{\text{ref}}^2/(2D_{\mu\mu})$  diagnose different failures of the constraint. Near  $\mu = 0$ , use a suitable distribution-level metric.

Compare these with an independently determined mixing time. A physically supported constrained interval requires

$$\tau_{\text{mix}} \ll t \ll \tau_\mu. \quad (100)$$

Choosing a small input pitch-angle rate establishes a synthetic asymptotic test, not this physical inequality. Adding two collision terms also requires a nonoverlapping process decomposition; it can otherwise double count encounters.

## 14 Independent homogeneous controls and encounters

### Lorentz pitch-angle evolution

For a fixed physical speed shell, the implemented operator is  $C_{\text{Lz}} = \nu \partial_\xi [(1 - \xi^2) \partial_\xi] / 2$ , with natural endpoint flux at  $\xi = \pm 1$ . The Legendre equation gives eigenrates  $\nu \ell(\ell + 1)/2$ . Gauss–Legendre projection and exact modal exponentials provide an independent solution; the tests additionally apply the differential operator to Legendre polynomials. Thus the  $\ell = 0, 1, 2, 3$  rates are  $0, \nu, 3\nu, 6\nu$ . Shell population and energy are preserved, while the  $\ell = 1$  momentum is transferred to the fixed reservoir. The coordinate map and full Jacobian are

$$u = v\xi, \quad \mu = \frac{mv^2(1 - \xi^2)}{2B}, \quad \left| \frac{\partial(u, \mu)}{\partial(v, \xi)} \right| = \frac{mv^2}{B}, \quad d^3v = \frac{2\pi B}{m} du d\mu. \quad (101)$$

The shell comparison changes  $\mu$  while preserving speed, as ordinary pitch scattering must; it is not an unbalanced bare moment diffusion term. All four evolved modes are also checked on actual mapped  $(u, \mu)$  nodes with the full transformed measure, rather than only checking the Jacobian at one point. Take  $F(v) = 1 + 0.3v^2$  on  $0.4 \leq v \leq 2$ ,  $m = 1.7$ ,  $B = 2.3$ ,  $\nu = 0.6$ ,  $t = 0.8$ , and amplitude  $\epsilon = 0.15$ .

Twelve positive Gauss nodes in each of  $v$  and  $\xi$  are mapped to  $(u, \mu)$ ; each speed shell is evolved by the production Legendre solver. An independent polynomial primitive  $I_p = \int_{0.4}^2 v^p F(v) dv$  gives

$$N = 4\pi I_2(1 + \epsilon\delta_{\ell 0}), \quad E = 2\pi m I_4(1 + \epsilon\delta_{\ell 0}), \quad P_{\parallel}(t) = \delta_{\ell 1} \frac{4\pi m \epsilon}{3} I_3 e^{-\nu t}.$$

The mapped quadrature reproduces count and energy to  $2 \times 10^{-14}$  relative test tolerance and the parallel momentum to  $2 \times 10^{-13}$  absolute tolerance. The  $\ell = 1$  momentum decreases, while all represented densities remain positive. This is a mapped shell quadrature check, not a rectangular  $(u, \mu)$  Lorentz discretization.

## Nonlinear conserving Dougherty evolution

The physical three-velocity operator is

$$C_D[F] = \nu \nabla_v \cdot [(\mathbf{v} - \mathbf{U})F + \theta \nabla_v F], \quad \mathbf{U} = \langle \mathbf{v} \rangle, \quad \theta = \frac{1}{3} \langle |\mathbf{v} - \mathbf{U}|^2 \rangle. \quad (102)$$

Integration against 1,  $\mathbf{v}$  and  $|\mathbf{v} - \mathbf{U}|^2$  proves that  $\mathbf{U}$  and  $\theta$  are constant for closed homogeneous evolution. For a positive normalized Gaussian mixture with weights  $p_a$ , component means  $\mathbf{U}_a$  and covariances  $A_a$ , the exact solution remains a positive mixture with

$$\mathbf{U}_a(t) = \mathbf{U} + e^{-\nu t}(\mathbf{U}_a(0) - \mathbf{U}), \quad A_a(t) = e^{-2\nu t} A_a(0) + \theta(1 - e^{-2\nu t})I. \quad (103)$$

Substituting each Gaussian into the now constant-moment Fokker–Planck equation verifies this solution; conservation makes it a solution of the original nonlinear model. It includes a non-Gaussian control, without substituting prescribed reservoir moments. Exact Gaussian gradients and Laplacians evaluate the strong operator independently of a centered temporal difference.

In `03_controls.py`, Lorentz modal errors are below  $6.5 \times 10^{-15}$ . The two-component Dougherty mixture has nonzero drift and anisotropic covariance. A  $40^3$  physical Cartesian velocity quadrature on extent 9 gives particle error  $2.0 \times 10^{-11}$ , covariance error  $7.82 \times 10^{-10}$  and centered-energy error  $4.76 \times 10^{-10}$ . The strong equation versus centered time derivative error is  $2.21 \times 10^{-10}$ . The sampled entropy increases; positive mixture reconstruction is analytic, and the smallest sampled density is  $1.96 \times 10^{-64}$ .

## Physical three-velocity Landau weak reference

The conventional one-species normalization used here is

$$C[F] = \Gamma \nabla_v \cdot \int F F' U(\mathbf{w}) [\nabla_v \ln F - \nabla_{v'} \ln F'] d^3 v', \quad U(\mathbf{w}) = \frac{I - \hat{\mathbf{w}} \hat{\mathbf{w}}^T}{|\mathbf{w}|}, \quad \mathbf{w} = \mathbf{v} - \mathbf{v}'. \quad (104)$$

With  $\int F d^3 v = n$ , the SI coefficient is  $\Gamma = q^4 \ln \Lambda / (8\pi \epsilon_0^2 m^2)$ , matching the mass convention of Fitzpatrick's Eqs. (3.93)–(3.94) [19]. Writing  $F = (n/v_*^3) \hat{F}$  gives the frequency scale  $\Gamma n / v_*^3$ . The Coulomb logarithm and its cutoffs are supplied parameters; this conversion does not calibrate the constrained scalar  $D$ .

For a normalized Gaussian of covariance  $A$ ,  $\mathbf{w}$  has covariance  $2A$  and the pair score is  $\mathbf{g} = -A^{-1}\mathbf{w}$ . The exact weak covariance and entropy rates are

$$\dot{A} = -\frac{\Gamma}{2} \mathbb{E}[\mathbf{w} \otimes U\mathbf{g} + U\mathbf{g} \otimes \mathbf{w}], \quad (105)$$

$$\dot{S} = \frac{\Gamma}{2} \mathbb{E}[\mathbf{g} \cdot U\mathbf{g}] \geq 0. \quad (106)$$

Trace conservation follows pointwise from  $U\mathbf{w} = 0$ . The center velocity integrates analytically; radial/polar/gyrophase quadrature retains the full physical  $d^3w$  geometry. A separate Cartesian Gauss–Hermite expectation avoids  $\mathbf{w} = 0$  by even quadrature order. A third independent check diagonalizes  $A$  and uses the Laplace integrals for  $1/r$  and  $1/r^3$ . Define  $Z(t) = \prod_j (1 + 4A_j t)^{-1/2}$  and  $V_j(t) = 2A_j/(1 + 4A_j t)$ ; then

$$\dot{A}_i = \frac{\Gamma}{\sqrt{\pi}} \int_0^\infty Z(t) \left[ t^{-1/2} \frac{V_i}{A_i} - 2t^{1/2} V_i \sum_j \frac{(1 + 2\delta_{ij})V_j}{A_j} \right] dt. \quad (107)$$

Gaussian second and fourth moments after multiplication by  $e^{-tr^2}$  give  $V_i$  and  $(1 + 2\delta_{ij})V_i V_j$ , proving the integrand. The substitution  $t = s^2$  removes its endpoint singularity before adaptive one-dimensional integration.

For  $A = \text{diag}(1.15, 1.15, 0.7)$  and  $\Gamma = 1$ , the covariance derivatives are

$$(-0.0740393, -0.0740393, 0.148079), \quad \dot{S} = 0.0413884.$$

The reported energy rate residual is  $2.1 \times 10^{-18}$ . Radial quadrature refinement reaches a  $4.14 \times 10^{-14}$  discrepancy from (107) at order 28, while the independent Cartesian Hermite route converges more slowly from 8.83% at order 8 to 0.316% at order 48. Tail and denominator-softening scans are saved separately. Softening changes only  $1/r$  and retains  $U\mathbf{w} = 0$ ; it is a convergence diagnostic, not an unmodified Coulomb tensor. A small-anisotropy check takes  $A = I + \epsilon \text{diag}(1, 1, -2)$ . At first order,  $\mathbb{E}|\mathbf{w}| = 4/\sqrt{\pi}$  for covariance  $2I$  and the angular moments are  $\mathbb{E}[n_i^2] = 1/3$  and  $\mathbb{E}[n_i^2 n_j^2] = (1 + 2\delta_{ij})/15$ . Substitution in the weak covariance rate gives  $\dot{A}_i = -4\Gamma \epsilon D_i / (5\sqrt{\pi}) + O(\epsilon^2)$ , where  $D = (1, 1, -2)$ . The instantaneous normalized stress rate approaches  $4\Gamma / (5\sqrt{\pi}) = 0.451351667$ ; for  $\epsilon = 0.03, 0.01, 0.003$  the two independent references agree and give 0.457605, 0.453335, 0.451936 respectively. This is a limiting conventional control, without a discovery claim.

Maxwellian score differences are parallel to  $\mathbf{w}$  and annihilate the kernel for multiple drifts and temperatures. A separate positive non-Gaussian nodal pair calculation checks weak number/momentum/energy conservation and positive entropy production. Its omitted coincident-node contribution is a stated quadrature prescription, not a general directional limit. These controls verify instantaneous weak moments; no full nonlinear Landau kinetic time integration or Gaussian closure trajectory is claimed.

## Direct screened encounter audit

`04_encounters.py` directly integrates equal-mass, equal-charge repulsive pairs in a uniform field with  $V(r) = \kappa e^{-r/\lambda}/r$ . It measures changes relative to the free helix at identical incoming conditions. Endpoint refinement holds fixed the incoming guiding center and phase at the  $z = 0$  crossing; changing the initial plane at fixed raw position and phase would instead change the orbit and is not a valid endpoint convergence test. For relative guiding-center offset  $b$ ,  $v_\parallel$  and gyrophase  $\varphi$  at that crossing, the incoming plane  $z = -L$  has phase  $\varphi + \Omega L/v_\parallel$  and transverse position  $(b - v_y/\Omega, v_x/\Omega)$ . This follows directly from the invariant free  $\mathbf{R} = \mathbf{r} + \mathbf{v} \times \mathbf{b}/\Omega$ .

A documented case uses  $b = 1.2$ ,  $v_\parallel = 1$ ,  $v_\perp = 0.35$ ,  $\Omega = 2$ ,  $\kappa = 0.05$ ,  $\lambda = 4$  and 32 phases. Increasing  $L$  from 12 to 32 changes the mean moment increment from  $5.0126 \times 10^{-5}$  to  $5.0074 \times 10^{-5}$  and its second moment from  $9.4135 \times 10^{-7}$  to  $9.4059 \times 10^{-7}$ . The individual incoming  $\mu$  is 0.00765625, so the rms increment is approximately 12.7% of that scale while the ensemble pair-energy error is below  $3.4 \times 10^{-11}$ . This is a controlled counterexample to inferring moment preservation solely from energy conservation; it does not characterize every weak encounter regime. For the same  $B = 2$  case, the maximum relative velocity change from the free helix is 0.0340, the minimum sampled separation is

1.045, and  $\rho/\lambda = 0.0219$ ,  $\rho/b_{90} \simeq 0.982$ ,  $\Omega\lambda/v_{\text{rel}} \simeq 7.55$ . Uniform  $B$  gives  $\rho/L_B = 0$ ;  $\nu/\Omega$  is unassigned because no encounter rate is measured. The free-orbit control has moment drift below  $1.2 \times 10^{-14}$  and guiding center drift below  $6 \times 10^{-16}$ . In a weak-field unscreened Rutherford control, increasing endpoint distance from 12 to 96 reduces the angular error, but the final relative discrepancy remains  $1.11 \times 10^{-3}$  and is recorded as finite-endpoint error. A separate close-encounter timestep scan decreases pair-energy error from  $5.40 \times 10^{-4}$  to  $3.57 \times 10^{-9}$ . The two individual moment increments are correlated by the imposed symmetric center-of-mass initialization; saved joint covariances must not be interpreted as independent species statistics.

No incident-flux/impact-parameter ensemble, plasma encounter frequency, or independently measured lifetime is assigned to these phase statistics.

## 15 Research questions and evidence

**Geometry and nullspace.** Begin with the proven finite-quadrature spectrum (33). Track its changes under a valid mirror, toroidal or dipole field. Separate collision null modes from decay generated through coupling to Hamiltonian transport. A small positive eigenvalue on one finite grid is not a continuum gap. Vary all resolutions, domains and velocity tails.

**Localization.** Compare the local kernel with a specified symmetric finite-range kernel, using a derived normalization as the range shrinks. Boundary renormalization must retain pair symmetry. A Gaussian-range surrogate remains a surrogate unless independent scattering physics supports it.

**Constrained dipole states.** Ref. [20] studies maximum-entropy dipole states and additional adiabatic constraints. Here the result sought is a dynamical relaxation law, an extra invariant, a loss of mixing, or an independently measured lifetime. Reproducing (38) is verification rather than discovery.

For any geometry-amplitude scaling, compare competing fits and inspect projector degeneracies before using regular perturbation theory. For steady-state parameter derivatives specify which particle, energy and marginal constraints are held fixed; otherwise a nullspace can make the derivative underdetermined.

### Executed prototype evidence and remaining limits

The public library has five substantive modules: **geometry**, **collisions**, **solver**, **controls** and the independent **reference**. Examples save inputs, boundary/model descriptions, package versions, source checksums and commit context with their results. All quoted values below are normalized, finite-grid evidence from the saved runs; rates are not dimensionally calibrated to a plasma.

The earlier 216-node NumPy oracle verifies the local uniform rate 0.11742603550295856 and its exact velocity-nullspace structure. The JAX implementation is compared with that independent pair calculation, a dense weak assembly, analytic Fourier decay, nonlinear directional linearization, conservation and discrete entropy identities. Geometry tests additionally cover both charges, both coordinate charts, sampled Jacobi and Liouville identities, independent finite-difference convergence, measure transformation, electric cross-field drift and mirror orbit timestep convergence.

**01\_relaxation.py** evolves a positive 105-node uniform nonlinear state for ten steps of  $\Delta t = 0.08$ . Its saved particle, energy and maximum moment-bin errors are below  $7.1 \times 10^{-17}$ ; the minimum step entropy increase is  $4.75 \times 10^{-4}$  and the maximum discrete entropy identity error is  $6.08 \times 10^{-15}$ . Nodal values remain positive, and the continuous spatial Fourier certificate and positive piecewise-linear velocity reconstruction are checked by the example. This is the specified nonlinear local surrogate, without self-consistent fields or a claim that its additional null directions disappear.

`02_geometry.py` evolves collision-only, ideal-only and combined linear states on the nonuniform vacuum toroidal annulus. On the base  $5 \times 3 \times 3 \times 11 \times 13$  grid at  $t = 0.24$ , the respective  $Q/Q(0)$  values are 0.820926, 1.000000 and 0.810980. Number, energy and moment-bin errors are below  $1.1 \times 10^{-17}$  for these three controls. Three-level scans separate radial, angular, parallel-velocity, moment and timestep errors; velocity tails are varied separately. The last changes in both final  $Q$  and dissipated fraction pass the 1% target for every scan. Changes in the dissipated fraction are 0.332% radially, rounding-level in the two angular grids for this resolved initial mode, 0.709% in parallel velocity, 0.638% in moment resolution, 0.495% in time, 0.0578% in parallel-velocity tail extent, and 0.519% in moment-tail extent. These are separate refinements of a specified observable, not a mathematical continuum convergence theorem. The unmodified 243-node collision matrix has 39 near-null eigenvalues, compared with the 39-dimensional independent invariant span from spatial populations and angular-weighted invariants in (53), with scaled eigenpair residual  $1.63 \times 10^{-15}$ . Every observed tiny-grid null direction is accounted for by that span; the eigenvalues themselves are unchanged. No continuum gap or uniqueness follows from that finite matrix.

`06_range.py` checks 27 combinations of spatial grid, velocity count and Gaussian range. The finest spatial density-rate discrepancy is at most  $3.03 \times 10^{-13}$  relative to (35). Underresolved narrow ranges are marked unresolved: for  $\ell = 0.15$ , 16 spatial nodes give a 58.5% density-rate error, while 32 give  $4.75 \times 10^{-4}$  and 64 reach rounding accuracy. This explicitly demonstrates why the range must be resolved before interpreting density-nullspace lifting.

`07_equilibria.py` independently integrates the stationary density (38) in all five fields with three velocity quadrature orders and two tail extents. The largest finest error on  $|u| \leq 7$ ,  $\mu \leq 20$  is  $1.61 \times 10^{-7}$  (dipole). Its separate full-marginal finite-quadrature energy solve gives  $\beta = 1.7310695$ , zero reported energy residual and moment-bin residual  $2.22 \times 10^{-16}$ . This verifies candidate construction, not attraction from arbitrary initial data.

`08_fields.py` performs actual nonlinear collision-only steps in mirror, dipole and perturbed-dipole boxes with the bounds derived above. At three spatial resolutions, equilibrium residuals are below  $2.8 \times 10^{-14}$  and initial entropy production is positive. The last spatial changes in initial production are 0.336%, 0.255% and 0.257% respectively. Three positive nonlinear steps on each coarse box conserve number, energy and the full moment marginal to below  $2.0 \times 10^{-17}$  in the saved run, with positive entropy increments. Velocity resolution/tails and longer nonlinear time evolution remain unresolved for these boxes. These runs impose natural collision flux only; they do not impose a confined dipole Hamiltonian boundary or demonstrate dipole self-organization.

Independent controls and direct encounters are described in the preceding sections. Performance measurements separate first-call compilation, warmed synchronization, pair chunk size and solver residual. Implicit parameter derivatives are checked against dense solves and centered finite differences. All performance claims remain specific to the saved device and accuracy; no unmeasured GPU speedup is claimed. The selected P1–P5 prototype examples have executable structural and convergence evidence. This does not certify every geometry or parameter regime; joint box time-evolution convergence remains specific to the executed scans. The continuous local zero set is classified above, including a critical-point failure of continuous extension; modified discrete energy maps require separate validation. Physical coefficient calibration, unequal-species closure, self-consistent electrostatics, and measured plasma mixing-versus-moment timescale separation are not implemented claims.

## Focused trajectory, unit and adjoint verification

On the symmetry axis of the harmonic mirror,  $B(z) = B_0 + az^2$ ,  $\dot{z} = u$  and  $m\dot{u} = -2\mu az$ . For initial  $z = 0$  and  $u = u_0 > 0$ , a separate scalar energy calculation locates the positive turning point from

$\mu B(z_t) = mu_0^2/2 + \mu B_0$  by bracketing. The bounce integral is

$$T_b = 4 \int_0^{z_t} \frac{dz}{\sqrt{2[E - \mu B(z)]/m}}.$$

The substitution  $z = z_t(1 - s^2)$  removes its endpoint singularity and gives  $T_b = 4 \int_0^1 2z_t [2\mu a z_t^2(2 - s^2)/m]^{-1/2} ds$ . Numerical quadrature is compared with turning-point events of a separately integrated five-coordinate Poisson flow. Both event position and period converge under integrator-tolerance refinement; the check uses nonunit mass and negative charge, and does not infer a collision rate from the orbit.

A uniform Fourier-pair regression also changes the base units rather than changing the physical input. Let the new units be  $L_*, T_*, M_*, Q_*, N_*$ , where  $N_*$  is a number-density unit. Then  $U_* = L_*/T_*$ ,  $E_* = M_* U_*^2$ ,  $B_* = M_*/(Q_* T_*)$ , and the prescribed uniform coefficient has unit  $D_* = M_*^2 L_*^2/(T_*^3 N_*)$ . In the  $(u, \eta)$  chart,

$$\hat{u} = u/U_*, \quad \hat{\eta} = \eta/E_*, \quad \hat{w}_v = w_v/(U_* E_*), \quad \hat{f} = f U_* E_*/N_*, \quad \hat{k} = k L_*.$$

Mass, charge, field and coefficient are divided by their respective units, and  $\partial_{\hat{u}} h = U_* \partial_u h$ . The independent pair calculation verifies  $\hat{C} = T_* U_* E_* C/N_*$  and  $\hat{\lambda} = T_* \lambda$  for the density-neutral observable  $\lambda = -\langle h, C \rangle_v / \langle f h^2 \rangle_v$ , for two unit maps and both charge signs. This check is restricted to the uniform reduced kernel; it does not replace the metric pullback required for a general coordinate or units change in phase space.

For the linear implicit system  $Ax = Mh$ ,  $A = M + \Delta t K$ , and observable  $J = x^T x$ , the adjoint obeys  $A^T \lambda = 2x$  and  $\partial_D J = -\Delta t \lambda^T K x$  at unit coefficient. The actual implicitly differentiated solver's VJP with respect to  $h$  returns  $M\lambda$ . Dividing by the positive diagonal mass recovers the adjoint, whose true residual is checked against the independently assembled dense matrix alongside the primal residual, dense derivative and centered finite-difference plateau. These are finite-grid derivative checks.

## Directional zero-set check

At the off-axis dipole state  $(1, 0.2, 0.3, 0.6, 0.4)$ , the independent `test_orbits.py` calculation approaches coincidence along  $u$  and  $\mu$  separately at separations  $10^{-2}, 10^{-3}, 10^{-4}$ . The two limiting  $PI_X P$  tensors remain different: there is no direction-independent continuous projector extension. For the smooth observable  $x^2 + u^3 + \mu x$ , however,  $\Delta \phi^T PI_X P \Delta \phi$  decreases quadratically along both paths. Thus identical-node diagonal contributions vanish in the weak form despite the kernel's directional discontinuity. This fact does not authorize deletion of distinct states with the same energy-flow vector; the general constructor rejects those ambiguous pairs. The uniform directional prescription is derived separately and tested against the exact oracle.

## Measured costs at matched error

For four independent uniform Fourier blocks with 128 velocity nodes each,  $\Delta t = 0.025$  and  $\lambda = 0.3$ , the common temporal error is  $2.77 \times 10^{-5}$ . On the Apple M4 CPU with JAX 0.11.2 and SOLVAX 0.28.2, five synchronized repetitions give the following measured values:

| Route        | Compile/first (s) | Median ( $\mu s$ ) | Range ( $\mu s$ ) | Temporary bytes |
|--------------|-------------------|--------------------|-------------------|-----------------|
| Dense        | 0.092             | 550.3              | 214.2–1297.8      | 262852          |
| PCG          | 0.184             | 571.7              | 273.1–804.7       | 16696           |
| Diagonal PCG | 0.225             | 43.8               | 34.5–126.6        | 16832           |

The implicit solve residuals are below  $6 \times 10^{-13}$ ; these separable uniform timings do not describe general pair-kernel complexity. For 945 actual weak pairs, a 16-pair chunk reduces XLA temporary buffers from 118913 to 16480 bytes, while the warm median increases from 96.3 to 158.2  $\mu$ s. The complete benchmark process peaks near 332 MB RSS, including the runtime and compiled executables. The base combined toroidal evolution takes 2.38 s for six steps after 1.12 s compilation/first execution and 1.22 s setup. Its state is 51480 bytes and its stored pair list contains 456885 unordered pairs; pair-application chunk workspace and all version/input/commit details are in the saved JSON. No extrapolated speedup or isolated-RSS interpretation is asserted.

## Nonuniform velocity quadrature and resolved tails

The additional `11_field_velocity.py` campaign evaluates the initial nonlinear entropy production on the same mirror, dipole and nonaxisymmetric boxes, with fixed  $3^3$  spatial quadrature. Its positive initial distribution is

$$f = \exp[-E - 0.2\mu + 0.1 \sin(x)u + 0.04y\mu]. \quad (108)$$

Positive Gauss–Legendre velocity quadrature replaces both the population weight  $w_X Bw_u w_\mu$  and the local pair weight  $w_X(Bw_u w_\mu)(Bw'_u w'_\mu)$  consistently. Derivatives and energy directions are unchanged. The three-point derivative is exact for the quadratic velocity energy and the linear- $u$  logarithmic perturbation on either node family. Number and energy are checked independently using shifted Gaussian integrals in  $u$  and exponential integrals in  $\mu$ .

The base velocity domain is  $u \in [-4, 4]$ ,  $\mu \in [0, 20]$ , with orders 25 and 21. Separate order scans are 17, 21, 25 in  $u$  and 13, 17, 21 in  $\mu$ . Tail extensions to  $|u| = 5$  and  $\mu = 24$  are separately checked at wider-domain orders 29 and 25. A Gauss tail extension changes node spacing; these extra order checks are essential and are not described as fixed-spacing tests. All 36 cases retain their inputs and status. Across all three fields, the largest final Gauss production change is 0.261%; the largest expanded-domain quadrature change is 0.180%. Base finite-domain number/energy quadrature errors are below  $9.4 \times 10^{-15}$ . Independent analytic omitted energy is about 0.050%, which is not itself a bound on production error.

Equidistant trapezoid controls at  $9^2, 17^2, 25^2$  velocity nodes remain unresolved: the last production changes are 16.63%, 10.14%, and 10.15% for mirror, dipole and nonaxisymmetric fields. Their statuses are preserved. The successful Gauss checks establish the stated initial-production comparison on a fixed spatial grid, not joint spatial–velocity or finite-time nonlinear convergence. The constructor now rejects zero pair directions without allocating and discarding a 25-component tensor per pair; the retained unit directions and collision action are unchanged and tested against explicit projector matrices.

## Finite-time campaigns and provenance

The historical discrete-gradient campaign and the current lagged-mobility campaign are separate time discretizations. The recovered four mirror rows retain their original execution at commit `fdc625d`; the subsequent field-split execution uses `4ce8e5c`. Dipole and nonaxisymmetric first steps in that execution failed their original GMRES check, and those failed logs and partial summaries remain archived. A failed Newton iteration is not a proof that a positive root does not exist.

The historical resume validator checked exact input hashes, package-source bytes and scientific-function syntax against its fixed anchor, together with recorded commits, clean-source hashes, environment metadata, complete time histories, positivity and retained solver/invariant tolerances. Tests retain that validator from Git history to audit the archived rows. Its saved summaries lack stage states and gradient-dependent entropy bounds, so rechecking them cannot replay every original entropy acceptance decision. Original timings and run identifiers remain attached to their producing executions.

The completed historical mirror run contains thirteen cases. Its eight independent refinements all pass the original 1% target; the largest selected-observable change is 0.1126103%, from extending  $|u| \leq 4$  to  $|u| \leq 5$ . The baseline relative entropy decreases by 43.7427% and the active moment by 25.0538% through  $T = 0.02$ . Across all accepted histories, number, energy and full moment-bin errors are at most  $7.779 \times 10^{-14}$ ,  $1.363 \times 10^{-13}$  and  $2.258 \times 10^{-14}$ ; the maximum nonlinear residual is  $7.545 \times 10^{-13}$  at target  $10^{-12}$ . The retained minimum nodal density is  $1.232 \times 10^{-16}$  and the largest absolute entropy identity discrepancy is  $5.877 \times 10^{-15}$ . These are fixed-other-parameter comparisons, not joint continuum convergence or a classification of every spatially localized constraint.

The archived validator independently accepts all thirteen rows and exactly reproduces each saved comparison. Raw summaries, environment, CPU report and logs are in `results/nonuniform_mirror_dg/raw`; its separate audit retains exact inputs and per-case costs. The  $7^3 \times 25 \times 21$  case has 180075 nodes and 41786320 bytes of stored grid arrays (not peak memory). Its four accepted steps took 2001.55 seconds including the first compilation, versus 366.68 seconds for the  $5^3$  case, on a four-logical-CPU AMD EPYC 7763 runner. These are observed run costs, not a matched-error speedup comparison.

The current `12_nonuniform_evolution.py` instead uses (88). Its fifteen predefined cases per field independently vary spatial order, parallel-velocity and magnetic-moment order, timestep, velocity tails and quadrature on each expanded domain. Baseline values are  $5^3 \times 25 \times 21$  nodes,  $D = 0.1$ ,  $T = 0.02$ , and  $\Delta t = 0.005$ ; the first-order timestep sequence is 0.005, 0.0025, 0.00125, 0.000625, 0.0003125. Every adjacent difference and the coarsest-to-finest difference are retained; a finest-pair pass does not certify the original baseline timestep or jointly refine the spatial and temporal discretizations. The common 1% observable-change target is retained. All units are normalized; these times use the freely prescribed coefficient.

`SM_EVOLUTION_FIELDS` and `SM_EVOLUTION_CASES` select predefined independent execution jobs, leaving their scientific inputs fixed. Without filters the script runs all forty-five cases. The CI driver separates each field/case pair and preserves failed or partial evidence. Historical discrete-gradient checkpoints are explicitly rejected by this new campaign. A completed selected job is distinct from a passing refinement comparison. Every row records its actual parameters, full step diagnostics, Newton history, separate KL contribution, residual entropy bound and original run identifier; metadata record source, versions and hardware. Independent refinements do not establish joint continuum convergence. No final outcome for the current campaign is asserted by this execution description.

**Completed mirror lagged-mobility campaign.** The retained collection contains 16 of 45 planned cases and remains unresolved overall. All fifteen independently executed mirror cases reach  $T = 0.02$  with the lagged-mobility entropy step, retaining a separate provenance record from the historical discrete-gradient campaign. The eight final, fixed-other-parameter comparisons all meet the original 1% target. Their largest change is 0.197278%, in the active-moment change for the final timestep pair. At  $\Delta t = 0.005$ , the  $5 \rightarrow 7$  spatial-order changes in entropy gain, final production and active-moment change are respectively 0.045351%, 0.024860% and 0.056944%. The corresponding largest tail change is 0.112113%.

The timestep sequence provides important qualification of these passes. For  $\Delta t = 0.005 \rightarrow 0.0025$ , entropy gain and active-moment change still differ by 1.30210% and 1.51359%. For the finest pair,  $0.000625 \rightarrow 0.0003125$ , the entropy-gain, final-production and active-moment discrepancies are 0.168910%, 0.131882% and 0.197278%. Successive absolute increments approximately halve, as expected for this first-order method. These are measured differences, not rigorous continuum error bounds; a first-order Richardson indicator additionally assumes the asymptotic expansion. The coarse-to-finest differences are still 2.45276%, 1.91519% and 2.85202%. A finest-pair pass therefore does not retrospectively certify the original coarse timestep. Spatial scans use the coarse timestep, whereas the

time scan fixes spatial order five; joint refinement is not asserted.

At  $\Delta t = 0.005$ , the lagged baseline relative entropy decreases by 42.5839%, compared with 43.7427% in the historical same-grid discrete-gradient baseline at the same timestep. The finest lagged result,  $\Delta t = 0.0003125$ , decreases relative entropy by 43.6546%.

The smallest stored historical discrete-gradient timestep is  $\Delta t = 0.0025$ ; its same-grid relative-entropy decrease is 43.7322%. Comparing this historical 0.0025 endpoint with the finest lagged 0.0003125 endpoint gives discrepancies of 0.177401% in entropy gain, 0.138914% in final production and 0.207378% in active-moment-change magnitude. Both comparisons hold all other grid, domain, initial-condition and coefficient inputs fixed. These cross-method differences do not identify either discrete endpoint with continuum truth. The separate numerical KL contribution to the entropy gain decreases from 3.51384% at the baseline timestep to 0.227342% at the finest; it is not included in the continuous nodal collision production.

Across the completed mirror histories, maximum number, energy and moment-bin errors normalized by the declared global scales are  $1.648 \times 10^{-12}$ ,  $4.312 \times 10^{-12}$  and  $4.862 \times 10^{-13}$ . Every stored entropy identity defect is within its recorded residual bound. These bounds certify step budgets, not spatial or temporal truncation error. Sampled meridional constraints improve with spatial order: the mirror  $\psi^2$  drift decreases from  $1.363 \times 10^{-10}$  at order three to  $1.594 \times 10^{-12}$  at order five and roundoff at order seven. This does not exhaust the disconnected Cartesian-cap population invariants or prove continuum mixing.

Accepted-step walls are observed costs of distinct shared Linux CI jobs, including each first-step compilation and excluding setup and separate initial/final diagnostics. The baseline four steps take 369.79 seconds; the finest timestep’s 64 steps take 1456.37 seconds. These costs are not a matched-error speedup comparison with the historical method. Stored grid array bytes exclude other workspace and line factors; process RSS is a high-water mark including compiled runtime. No dipole or nonaxisymmetric campaign conclusion is inferred from these mirror results.

The public mirror audit retains the original per-case provenance and recomputes all eight comparisons, including every adjacent timestep difference and the coarse-to-finest bias.

**A retained dipole solver failure.** At the same frozen campaign revision 3d842bee, the dipole  $5^3 \times 25 \times 13$  case with  $U = 4$ ,  $M = 20$  and  $\Delta t = 0.005$  accepts one step, then rejects a PCG correction. The native convergence status accompanies a recomputed relative residual  $3.118 \times 10^{-4}$ , above both the requested  $10^{-6}$  and the unchanged  $5 \times 10^{-6}$  acceptance guard. The accepted first step uses seven Newton corrections and 7104 total PCG iterations, preserves every recorded moment-bin population, and has minimum nodal density  $5.435 \times 10^{-232}$ . Its entropy increase  $6.9723 \times 10^{-5}$  splits into frozen-mobility production  $6.3797 \times 10^{-5}$  and numerical KL  $5.9261 \times 10^{-6}$ . The full correction fractions and large negative log increments are retained in the public aggregate’s partial-case record. The extreme evolved tails make finite-precision conditioning a concern; a positive exact discrete root does not guarantee convergence in binary64. This failure is not a physical decay measurement or a converged dipole refinement. No clipping, kernel change or relaxed tolerance is used to turn it into a passing result.

**Solver-tolerance audit on one fixed mirror grid.** Example 19 at source revision e1bc1ab retains every state in five four-step runs on  $3^3 \times 25 \times 21$  nodes, with  $D = 0.1$ ,  $T = 0.02$  and  $\Delta t = 0.005$ . Outer Newton tolerances  $10^{-10}$ ,  $10^{-11}$ ,  $10^{-12}$  at inner tolerance  $10^{-6}$  give byte-identical endpoints; this is not an outer-error-order estimate. At outer tolerance  $10^{-12}$ , the separate inner sequence is  $10^{-4}$ ,  $10^{-5}$ ,  $10^{-6}$ . An independent NumPy saved-state audit reconstructs all 20 nonlinear residuals and entropy/invariant budgets, with explicit unordered-pair production checks. The maximum independently recomputed

nonlinear relative residual is  $2.162 \times 10^{-13}$ ; maximum cumulative number, energy and full moment-bin errors are respectively  $1.019 \times 10^{-14}$ ,  $1.579 \times 10^{-14}$  and  $4.135 \times 10^{-15}$ . The last inner pair has entropy-scaled log endpoint difference  $1.393 \times 10^{-14}$  and maximum independently recomputed relative change in the three principal observables  $5.27 \times 10^{-13}$ , well below the unchanged  $10^{-8}$  endpoint and  $10^{-6}$  observable targets. The fixed finite-dimensional reference is not continuum truth. Krylov correction vectors are not archived, so their true residual scalars are checked as producer records rather than independently replayed vectors. The five recorded run costs sum to 89.698 seconds on the Apple M4; the separate 3.670-second setup/initial-diagnostic cost is excluded from that sum. These observed costs retain first-use compilation, evidence writes and concurrent load, and supply no speedup claim. Exact normalized inputs, versions and hardware are retained in the producing summary.

## Bounded Maxwellian encounter flux and thermal center motion

Let  $\theta = k_B T/m$  denote a single-particle velocity variance. For two equal Maxwellian particles,  $\mathbf{w} = \mathbf{v}_1 - \mathbf{v}_2$  and  $\mathbf{V} = (\mathbf{v}_1 + \mathbf{v}_2)/2$  are independent, with covariances  $2\theta I$  and  $\theta I/2$ . The transformation has absolute Jacobian one. The relative distribution and equations for a screened pair potential are

$$F_{\text{rel}}(\mathbf{w}) = (4\pi\theta)^{-3/2} e^{-|\mathbf{w}|^2/(4\theta)}, \quad (109)$$

$$\dot{\mathbf{r}} = \mathbf{w}, \quad \dot{\mathbf{w}} = \Omega \mathbf{w} \times \mathbf{e}_z + \frac{2\kappa e^{-r/L_s} (1 + r/L_s)}{m r^3} \mathbf{r}, \quad \Omega = qB/m. \quad (110)$$

The conserved relative energy is  $m|\mathbf{w}|^2/4 + \kappa e^{-r/L_s}/r$ ; center motion follows an independent free helix. The new six-dimensional relative integrator is checked against the independent twelve-dimensional pair integrator, including nonunit masses and both charge signs.

For guiding-center impact  $b$ , the translation to the incoming perpendicular position has unit Jacobian. Axisymmetry integrates its orientation to  $2\pi b db$ . On a plane normal to the field, the two-sided incoming crossing measure per partner density is

$$\frac{d\nu}{n} = 2(2\pi b db) w_{\parallel} F_{\text{rel}}(\mathbf{w}) w_{\perp} dw_{\perp} dw_{\parallel} d\varphi, \quad w_{\parallel} > 0. \quad (111)$$

The factor two includes the opposite incoming direction. There is no unordered-pair half factor for this tagged-particle flux, and the plane flux uses  $|w_{\parallel}|$ , not the cross-section flux  $|\mathbf{w}|$ . For impact area  $A = \pi(b_{\text{max}}^2 - b_{\text{min}}^2)$ , the bounded integral is  $2A(2\pi)J_{\parallel}J_{\perp}/(4\pi\theta)^{3/2}$ , where  $J(a, b) = 2\theta[e^{-a^2/(4\theta)} - e^{-b^2/(4\theta)}]$ . This supplies an independent analytic normalization.

Undo outgoing free gyromotion when defining  $\Delta\mathbf{w}_{\perp} = \overline{\mathbf{w}}_{\text{out},\perp} - \mathbf{w}_{\text{in},\perp}$ . With  $c = m\Delta|\mathbf{w}_{\perp}|^2/(8B)$ , the individual changes are  $\Delta\mu_{1,2} = c \pm [m/(2B)]\mathbf{V}_{\perp} \cdot \Delta\mathbf{w}_{\perp}$ . Exact Gaussian center averaging therefore gives

$$\langle \Delta\mu_1 \rangle_V = c, \quad (112)$$

$$\langle (\Delta\mu_1)^2 \rangle_V = c^2 + \sigma_c^2, \quad \langle \Delta\mu_1 \Delta\mu_2 \rangle_V = c^2 - \sigma_c^2, \quad (113)$$

$$\sigma_c^2 = \frac{m^2\theta}{8B^2} |\Delta\mathbf{w}_{\perp}|^2. \quad (114)$$

A separate Gaussian quadrature of the actual individual initial/final velocities verifies these raw moments. A zero-center trajectory omits  $\sigma_c^2$ . The drift and diffusion moment integrals are  $A_{\mu}/n = \int \langle \Delta\mu \rangle_V d\nu/n$  and  $D_{\mu\mu}/n = \frac{1}{2} \int \langle (\Delta\mu)^2 \rangle_V d\nu/n$ ; the latter half is the second-moment convention, not pair counting.

The executed normalized case uses  $m = q = 1$ ,  $B = 2$ ,  $\kappa = 0.03$ ,  $L_s = 3$ ,  $\theta = 0.5$ ,  $b \in [1.2, 2.4]$ ,  $w_{\parallel} \in [0.7, 1.5]$  and  $w_{\perp} \in [0.2, 0.8]$ . The  $4^3 \times 16$  base quadrature is independently refined in each impact/speed coordinate through orders 3, 4, 5, in gyrophase through 8, 16, 32, and by a half-grid

phase shift. Moving the planes through  $L = 18, 24, 32$  preserves the same incoming free helix. The first outward exit is recorded; reflected trajectories are retained, and failure to exit within  $6L/w_{\parallel}$  raises. Eight held-out states additionally check active maximum steps 0.2, 0.1, 0.05. Direct quadrature changes are below 0.021%; analytic flux normalization error is  $1.62 \times 10^{-9}$ . The bounded results are  $\nu/n = 1.2601061$ ,  $A_{\mu}/n = 3.11546 \times 10^{-6}$  and  $D_{\mu\mu}/n = 8.00765 \times 10^{-7}$  in the stated normalized units.

A supplementary multilinear scattering table is tested on 64 uniformly sampled states in the inner quadrature-node hull (seed 701). Its held-out second-moment relative RMS error is 20.5%, exceeding the predeclared 10% target; this surrogate remains unresolved despite converged direct integrals. These speed bands contain only 11.6% of the full thermal incoming flux in the same annulus. Omitted impact/speed regions and successive encounter correlations remain untested. Truncation can exclude reverse encounters, so the positive conditional drift is not a full Maxwellian-plasma drift. Neither these bounded crossing coefficients nor the interpolation checks calibrate the constrained scalar  $D$  or establish a dipole lifetime.

## Broader encounter coverage and microscopic reversal

Example 13 retains the same screened pair model and thermal center average, while widening the tested speed band to  $w_{\parallel} \in [0.35, 3]$ ,  $w_{\perp} \in [0, 2.5]$  and splitting the impact interval into  $[0.15, 0.4]$ ,  $[0.4, 0.8]$ ,  $[0.8, 1.2]$ ,  $[1.2, 2.4]$ ,  $[2.4, 4.8]$  and  $[4.8, 8]$ . The impact variable  $b^2$  is sampled uniformly; each squared speed is drawn from its truncated exponential incoming-flux law. These inverse cumulative distributions implement the positive crossing measure above. Scrambled Sobol orders 32, 64, 128, 256 share one scramble; a separate 128-point scramble checks replication. Gyrophase uses 16 equally spaced points and an independent  $8 \rightarrow 16$  refinement. The two-scramble spread is a finite numerical comparison, not a statistical confidence interval. The speed band retains 88.8641% of the complete thermal crossing flux in the same impact annuli; this flux fraction does not bound the omitted moment contribution.

Microscopic reversal is checked on 16 separate transmitted encounters. Starting at the actual outgoing position, apply parity and time reversal,  $(\mathbf{r}_R, \mathbf{w}_R) = (-\mathbf{r}_{\text{out}}, \mathbf{w}_{\text{out}})$ , and reverse the charge, hence  $\Omega_R = -\Omega$ . A common axial rotation restores the incoming guiding-center orientation, and the free-helix phase is recovered from that state. The relative equations then retrace the encounter. The thermal-center mean changes sign, while its second and cross moments are unchanged. Maximum absolute discrepancies are respectively  $2.87 \times 10^{-11}$ ,  $1.09 \times 10^{-13}$  and  $7.52 \times 10^{-14}$ . This tests the two-body inverse map; it does not establish detailed balance for a truncated incident ensemble that omits reverse states.

The broad campaign evaluates 38,220 distinct trajectories. All base trajectories exit, but one independent-replica state fails to reach an outgoing plane within  $6L/w_{\parallel}$ :

$$(b, w_{\parallel}, w_{\perp}, \varphi) = (0.9938225266, 1.9056642462, 1.9668376625, \pi/2), \quad L = 24. \quad (115)$$

Missing samples are not renormalized away. The  $[0.8, 1.2]$  annulus and total conditional drift/diffusion are therefore reported as unresolved, with no aggregate coefficient. The final empirical diffusion spreads in the six annuli are 10.4%, 11.6%, 65.9%, 8.62%, 3.95% and 4.85%. Only the outer two meet the stated 5% finite-comparison target; the inner four do not. Positive contributions from completed nodes are lower bounds only on their specified finite positive quadrature sums, not certified lower bounds on a continuous integral.

Even in the previously tested  $[1.2, 2.4]$  impact annulus, the wider speed band gives  $D_{\mu\mu}/n = 2.78117 \times 10^{-5}$ , about 34.7 times the earlier narrow-band value, and reverses the conditional drift sign to  $A_{\mu}/n = -1.01609 \times 10^{-5}$ . This row is itself unresolved by the 8.62% spread and 5.26% phase change; the comparison demonstrates cutoff sensitivity rather than a converged full rate. Twelve separate states

give relative  $L^2$  changes of the twelve-state second-moment vector of  $4.57 \times 10^{-5}$  on moving the planes from 24 to 32 and  $1.35 \times 10^{-10}$  on decreasing the active maximum timestep from 0.1 to 0.05. These state checks do not settle the remaining ensemble quadrature error or its failed trajectory.

For a repulsive uniform-field pair,

$$\ddot{z} = k(t)z, \quad k(t) = \frac{2\kappa e^{-r/L_s}(1 + r/L_s)}{m r^3} \geq 0. \quad (116)$$

On  $z < 0$ , the incoming positive parallel velocity decreases monotonically. If it turns, the particle thereafter recedes toward the incoming exit and cannot turn again. If it crosses  $z = 0$  with positive velocity, that velocity increases on  $z > 0$ . Reaching  $(z, \dot{z}) = (0, 0)$  at a finite time is excluded by uniqueness of the regular equations unless the trajectory stays in that plane. Thus an exiting trajectory has at most one parallel turn. A finite flight-budget failure can involve slow outward travel or long near-turning motion; it does not establish a bound pair or a failure of molecular chaos. The omitted impact and speed tails, reverse-complete incident statistics and successive-encounter correlations remain unchecked. No physical constrained coefficient, plasma relaxation time or dipole lifetime follows.

## Resolving the failed encounter's flight budget

Example 17 preserves the original broad-campaign metadata and replays its exact failed node. The declared flight factors 6, 12, 24, 48 vary only the finite integration budget. Factor 6 ends at  $t = 75.5642$ , after reflection, with  $z = -19.268$  and  $\dot{z} = -0.320016$ ; the particle is already travelling toward the incoming plane. Factor 12 reaches the reflected outgoing event  $z = -24$  at  $t = 90.3503$ , with one parallel turn. Factors 24 and 48 reach the same event. This diagnoses an insufficient original numerical flight budget, not a trapped trajectory.

Active maximum steps 0.2, 0.1, 0.05 give a final conditional second-moment change  $2.13 \times 10^{-10}$ ; independently moving the planes through  $L = 24, 32, 40$  gives a final change  $2.24 \times 10^{-5}$ . The original comparison of the 0.2 relative solve to the Cartesian pair path has velocity error  $1.69 \times 10^{-8}$ , above the unchanged  $10^{-8}$  absolute target, and is retained as an unresolved coarse comparison. The finest declared relative solve instead uses maximum step 0.05 and its own event time  $t = 90.350330908$ . Independent twelve-dimensional pair paths with maximum steps 0.05 and 0.025 agree with it in relative velocity to  $5.15 \times 10^{-12}$  and  $4.34 \times 10^{-11}$ , and in relative position to  $2.54 \times 10^{-10}$  and  $4.02 \times 10^{-9}$ . Their velocity plateau is  $4.63 \times 10^{-11}$ . Both satisfy the unchanged  $10^{-8}$  position and velocity targets. No discrepancy is removed by an invariant correction.

Restoring this one node changes only the chosen replica's finite quadrature. Its conditional thermal-center mean, second and cross moments are  $(0.2205711, 0.1127298, -0.0154266)$ . The positive increment to that replica's finite  $D_{\mu\mu}/n$  sum, using the half-second-moment convention, is  $4.90439 \times 10^{-5}$ , equal to about 20.12% of the *base* annulus's reported finite diffusion sum, even though the node represents only  $1/2048 = 0.04883\%$  of the replica count. The two populations are distinct; this ratio is a diagnostic of numerical leverage, not a percentage correction to a converged continuous coefficient. For this node the unperturbed relative gyro-radius is  $\rho_{\text{rel}} = w_{\perp}/|\Omega| = 0.9834188$ . At the declared zero-plane phase  $\varphi = \pi/2$  and  $\Omega = 2$ ,  $\mathbf{r}_{\perp} = (b - w_{\perp} \sin \varphi/\Omega, w_{\perp} \cos \varphi/\Omega)$ ; the free relative orbit's closest perpendicular gap is only  $b - \rho_{\text{rel}} = 0.0104037$ . The interacting independent Cartesian path has sampled minimum separation 0.0225124. These are relative-orbit geometrical quantities, not individual thermal gyro-radii. The Coulomb reference length is  $b_{90} = \kappa/[(m/2)|\mathbf{w}_{\text{in}}|^2] = 0.00799999$ , so the free gap is  $1.30046b_{90}$ , despite the much larger guiding-center impact. This is the unscreened Coulomb 90-degree reference length, not an exact deflection parameter for the screened magnetized trajectory. At the sampled closest point, the actual relative perpendicular velocity gives  $\rho_{\text{kin}} = |\mathbf{w}_{\perp}(t)|/|\Omega| = 0.683913$  and  $\rho_{\text{kin}}/r = 30.3794$ ; these differ from the incoming ratio  $\rho_{\text{rel}}/r = 43.6835$ . The measured full-speed ratio

is  $|\Omega|r/|\mathbf{w}(t)| = 0.0302907$ , while  $|\Omega|b_{90}/|\mathbf{w}_{\text{in}}| = 0.00584237$ . These local scales place this concrete close encounter outside a verified grazing or close-collision adiabatic ordering. They help interpret its large finite moment contribution, but are not a counterexample to an asymptotic SM working regime or an ensemble relaxation law. Resolving the trajectory does not settle the independent Sobol/phase spread, omitted bands or kinetic closure. Example 13’s original aggregate remains unresolved and its metadata and failed comparisons remain available.

The same final audit gives a deterministic inverse-cutoff counterexample, with the field explicitly fixed at  $B = 2$  in both trajectory and thermal-center moment calculations. The incoming node  $(b, w_{\parallel}, w_{\perp}, \varphi) = (1.2, 0.7, 0.8, \pi/2)$  maps under the tested microscopic inverse to  $(1.20082758, 0.69429998, 0.80495188, 3.94743171)$ . Both inverse speeds lie outside the original narrow bands. Its forward mean and second moment are  $4.96721 \times 10^{-4}$  and  $7.44799 \times 10^{-6}$ . The reverse mean negation, second-moment equality and cross-moment equality have absolute errors respectively  $5.07 \times 10^{-16}$ ,  $3.19 \times 10^{-18}$  and  $2.18 \times 10^{-18}$ ; energy errors are below  $1.52 \times 10^{-14}$ . This is a fixed boundary counterexample, not an independent statistical validation or a fitted state. The outgoing event is transverse, so continuity also gives nearby interior incoming states with inverse speeds outside the bands. The incident rectangle therefore is genuinely not closed under this inverse. Conditional drift on that rectangle is not a full Maxwellian drift or evidence against microscopic reversibility.

### Positive periodic scattering table: bulk accuracy and a failed tail target

The refined table interpolates the signed conditional mean  $c$  and  $\ln \sigma_c^2$  with cubic interpolation in impact and speeds and a truly periodic cubic spline in gyrophase. Reconstruction as  $c^2 + \sigma_c^2$  and  $c^2 - \sigma_c^2$  preserves the positive variance and thermal-center identities. A  $4^3 \times 16$  version, tested at 256 fresh uniform states with seed 1301, gives a second-moment relative  $L^2$  error 1.038%, but a 95th percentile pointwise relative error 9.869%; the prescribed 5% tail target fails.

Example 16 refines to  $6^3 \times 32 = 6912$  training states on the original narrow bands, and evaluates 256 independent uniform states (seed 1311) inside the training-node hull, with no extrapolation. Its relative  $L^2$  error is defined as

$$e_{\text{bulk}} = \frac{\sqrt{\sum_i (\widehat{M}_i - M_i)^2}}{\sqrt{\sum_i M_i^2}}, \quad e_i = |\widehat{M}_i - M_i|/M_i, \quad (117)$$

where  $M_i = \langle (\Delta\mu)^2 \rangle_V > 0$  is the direct moment. This first sample gives  $e_{\text{bulk}} = 0.1328\%$  and an observed 95th percentile 4.480%, meeting the original observed-error criteria. However, the two-sided 95% binomial order-statistic interval for the population 95th percentile is [2.273%, 30.303%]; 256 states do not establish a 5% population tail bound. Eight independent impact-direction derivative checks give relative  $L^2$  error 0.3632%, with direct centered-difference plateau change  $6.73 \times 10^{-7}$ . This is one moment derivative on eight states, not a general kinetic sensitivity result.

The same table is then tested without retraining on 4096 new independent uniform states (seed 1321). All direct trajectories exit. The original 256-state measurements and training archive remain separate; cached nodes, producer metadata, dependency hashes and interpolant values are verified. After a duration-helper change invalidated that cache’s controls hash, the complete training table, original 256-state stage and fresh 4096-state stage were regenerated at source revision `cdb577b`. All eleven saved training, moment, derivative and validation arrays are identical in shape, dtype and element bytes to the archived `90abe1d` measurements; the negative tail result is reproduced without retuning. The fresh bulk error is 0.2282%, with an empirical 2000-resample bootstrap interval [0.1912%, 0.2700%]. The observed 95th percentile is 6.2610% and its two-sided 95% order-statistic interval is [5.3207%, 7.2237%]. Thus the predeclared 5% tail target fails in the larger sample, including its confidence interval. Of these

states, 5.884% have relative error above 5% and 0.2930% above 100%; small direct moments permit large relative outliers despite small bulk error.

The order-statistic interval uses the binomial distribution of the count below a fixed population quantile, valid for these independent uniform samples and a continuous error distribution. It describes that declared sampling distribution, not every state, a flux-weighted ensemble or omitted domain regions. The bootstrap bulk interval is an empirical uncertainty estimate. A failed tail target is retained as a result, with no parameter retuning on those validation states. Table resolution and conditional moment accuracy do not supply kinetic closure, screening validation, inter-encounter independence or a constrained collision coefficient.

## Predeclared eight-node refinement on the unchanged validation domain

A separate, predeclared  $8^3 \times 32$  table uses 16384 direct training states, with the same signed-mean/log-variance interpolation, physical inputs, integration tolerances and accuracy targets. The validation domain remains exactly the original six-node inner hull: the impact and two speed intervals are approximately  $[1.2405183, 2.3594817]$ ,  $[0.7270122, 1.4729878]$  and  $[0.2202591, 0.7797409]$ , respectively, with uniform gyrophase on  $[0, 2\pi)$ . The larger eight-node training hull is not substituted for this fixed test domain. Gauss-six and Gauss-eight training nodes have no bit-identical full states, so no earlier trajectory is reused. The failed six-node record above remains available and is not revised by the new outcome.

Source revision `e1bc1ab` declares 4096 fresh primary states with seed 1331 and bootstrap seed 1332, and 256 fresh secondary states with seed 1333 and bootstrap seed 1334, before training or inspecting their outcomes. The normalized inputs remain  $m = q = 1$ ,  $B = 2$ ,  $\kappa = 0.03$ ,  $\lambda_s = 3$ ,  $\theta = 0.5$ , incoming/outgoing planes  $L = 24$ , maximum DOP853 step 0.2 and relative tolerance  $10^{-10}$ . The exact hull, source, versions and hardware are recorded in the public normalized-input metadata.

All 4096 primary direct encounters exit. Their second-moment bulk relative error is 0.11506%, with 2000-resample bootstrap interval  $[0.08909\%, 0.14119\%]$ . The observed 95th percentile is 2.53760%; the two-sided 95% order-statistic interval is  $[2.05643\%, 3.36706\%]$ . Thus the unchanged criteria—bulk error at most 2%, observed 95th percentile at most 5%, and confidence upper limit at most 5%—all pass for this declared sampling distribution. The secondary sample has bulk error 0.05963% and observed 95th percentile 3.90539%, but its much wider interval  $[1.25255\%, 10.00161\%]$  illustrates why the larger primary sample is the confidence evidence. The eight impact-direction checks give derivative relative error 0.47585%, below the original 5% target; their direct finite-difference plateau change is  $6.7149 \times 10^{-7}$ . Maximum saved training, secondary and primary energy errors are respectively  $5.76 \times 10^{-10}$ ,  $3.59 \times 10^{-10}$  and  $4.39 \times 10^{-10}$ , below the unchanged  $10^{-8}$  guard applied jointly to all three stages.

The pass is not a pointwise accuracy bound: 140 of 4096 states (3.4180%) exceed 5% relative error, and eight (0.1953%) exceed 100%; the maximum relative error is 293.95%. Small direct moments still permit large relative outliers. An independent saved-array audit reconstructs the seeds and fixed domain, verifies no overlap with earlier states, rebuilds the interpolant and reproduces the reported statistics and derivative comparisons. It verifies source dependencies against the producer revision. Training energy maxima are replayed from the saved checkpoint; secondary and primary energy checks use saved producer maxima rather than rerunning individual trajectories.

The calculation comprises  $16384 + 4096 + 256 + 48 = 20784$  direct solves, including 48 finite-difference solves for the eight derivative states, and takes 354.638 seconds on the recorded Apple M4. This is one observed whole-run cost including concurrent machine load, not a matched-cost speedup over the older table. The successful bounded conditional interpolation test does not validate a physical kinetic closure,

a calibrated collision coefficient, inter-encounter independence or accuracy outside the declared hull.

## Reproducible evolution visualization

The README animation is a solution of the uniform linear surrogate at finite velocity quadrature, not a nonlinear encounter or physical-rate simulation. The nine positive velocity masses  $M_j$  are proportional to  $\exp(-v_j^2/2)$ , with  $v_j$  equally spaced on  $[-2, 2]$  and  $\sum_j M_j = 1$ . Define

$$p_j = \frac{v_j^2 - \sum_i M_i v_i^2}{\max_i |v_i^2 - \sum_l M_l v_l^2|}, \quad h_j(x, 0) = (0.25 + 0.20p_j) \cos x. \quad (118)$$

The density and velocity-neutral amplitudes are evolved using matrix exponentials of an independently assembled ordered-pair Fourier matrix ( $N_x = 64$ , Gaussian range  $\ell = 0.6$ ) and the analytic local matrix. For  $D = 0.8$ ,  $B = 1.7$ ,  $q = -1.3$ , and  $\tau = \lambda t$ , the exact comparisons are

$$h_{j,\text{local}}(x, t) = [0.25 + 0.20e^{-\tau} p_j] \cos x, \quad (119)$$

$$h_{j,\text{range}}(x, t) = [0.25e^{-(1-e^{-0.6^2/2})\tau} + 0.20e^{-\tau} p_j] \cos x. \quad (120)$$

All 65 saved frames on  $0 \leq \tau \leq 8$  are checked against these formulas;  $1 + h_j$  stays positive. Fixed color limits prevent apparent relaxation caused by automatic rescaling. The displayed velocity coordinate is a discrete node index, not a physical three-dimensional velocity grid. GIF and MP4 use the same frames, and the arrays are saved in a compressed archive. The accompanying validation and cost panels read the existing experiment JSON files. Their metadata retain each source file's hash and source commit, rather than attributing old measurements to the visualization revision.

## Continuous local geometric sensitivity

Example 14 differentiates an initial local entropy-production density, not a kinetic trajectory or an implicitly defined steady state. Its prescribed distribution is

$$\ln f = -E - 0.2\mu + p, \quad p = 0.1 \sin(x)u + 0.04y\mu. \quad (121)$$

At  $\mathbf{X} = (1, 0.12, 0.3)$ ,  $m = 1.3$ ,  $q = -0.8$ ,  $D = 0.1$ , define

$$\mathcal{P}(\mathbf{X}) = \frac{DB^2}{2} \iint f f' |I_X P \Delta A \ln f|^2 du d\mu du' d\mu'. \quad (122)$$

Energy and moment degeneracy give  $P \Delta A \ln f = P \Delta A p$ . The independent NumPy reference reconstructs  $A p$  from the analytic gradient  $\nabla_X p = (0.1 \cos(x)u, 0.04\mu, 0)$ ,  $p_u = 0.1 \sin(x)$  and the vacuum flow (58); it shares only the field derivative data. Both evaluations retain the two local  $B$  Jacobians. Automatic differentiation translates the sample in  $x$  while holding the analytic distribution definition, velocity bounds and field parameters fixed; hence it differentiates  $B^2$ ,  $f$ ,  $f'$ , both observable actions and the energy-flow projector. No part of the geometry is stopped in this derivative.

On the base  $25 \times 21$  velocity quadrature,  $|u| \leq 4$ ,  $0 \leq \mu \leq 20$ , the mirror, dipole and perturbed dipole respectively give

$$(\mathcal{P}, \partial_x \mathcal{P})_m = (0.00142987874, -0.00316286298), \quad (123)$$

$$(\mathcal{P}, \partial_x \mathcal{P})_d = (0.00528569785, 0.0324274535), \quad (124)$$

$$(\mathcal{P}, \partial_x \mathcal{P})_{na} = (0.00507509906, 0.0310717576). \quad (125)$$

The independent pair values differ by at most  $1.03 \times 10^{-15}$  relative. Sixteen independent centered-difference steps from  $0.04$  to  $10^{-8}$  show the derivative plateau; the last-three-step maximum relative discrepancy is below  $1.54 \times 10^{-8}$ . Separate velocity-order checks, expanded  $|u| \leq 5$  and  $\mu \leq 24$  bounds, and quadrature refinement on each expanded domain give final derivative changes below  $0.661\%$ . The local maps are injective on these domains, as established above. This is a checked derivative of the specified normalized weak observable, with no general steady-state sensitivity or physical-rate inference.

## Positive dipole distributions with incompatible flux populations

Example 15 constructs a finite-domain counterexample without a time solve. Use poloidal flux  $\psi$ , magnetic latitude  $\lambda$ , and periodic azimuth  $\theta$ . With  $r = C \cos^2 \lambda / \psi$ ,

$$\left| \frac{\partial \mathbf{X}}{\partial(\psi, \lambda, \theta)} \right| = \frac{C^3 \cos^7 \lambda}{\psi^4}, \quad B d\mathbf{X} = \frac{C \cos \lambda \sqrt{1 + 3 \sin^2 \lambda}}{\psi} d\psi d\lambda d\theta. \quad (126)$$

An independent Cartesian automatic-differentiation determinant checks both identities at two nontrivial coordinates and verifies the inverse flux map. The finite domain is  $0.4 \leq \psi \leq 1$ ,  $|\lambda| \leq \pi/6$ ,  $|u| \leq 4$ ,  $0 \leq \mu \leq 12$ ;  $C = 1$  and  $f_* = \exp(-E - 0.2\mu)$ . Angular integration is exact over  $2\pi$ .

Let conditional means under  $f_* d\Gamma$  at fixed  $\mu$  be denoted by an overbar, and set

$$r_E = E - \bar{E}_\mu, \quad r_\psi = \psi - \bar{\psi}_\mu, \quad r_2 = \psi^2 - \bar{\psi}_\mu^2. \quad (127)$$

Solve the positive two-dimensional covariance system

$$\int f_* \begin{pmatrix} r_E^2 & r_E r_\psi \\ r_E r_\psi & r_\psi^2 \end{pmatrix} d\Gamma \begin{pmatrix} a \\ b \end{pmatrix} = \int f_* \begin{pmatrix} r_E r_2 \\ r_\psi r_2 \end{pmatrix} d\Gamma, \quad h = r_2 - a r_E - b r_\psi. \quad (128)$$

The Gram matrix is nonsingular on this positive product domain: an affine relation between  $E$  and  $\psi$  would require a vanishing  $u^2$  coefficient, then a constant flux. The residual  $h$  is nonzero because  $\psi^2$  cannot be an affine function of  $\psi$  on a nonzero-width flux interval after the  $u^2$  coefficient has vanished. Define  $f_\pm = f_*(1 \pm \epsilon h)$ . The conservative bound

$$|h| \leq (\psi_{\max}^2 - \psi_{\min}^2) + |a| E_{\max} + |b| (\psi_{\max} - \psi_{\min}) =: H \quad (129)$$

holds over the entire truncated domain, since conditional means lie in the same global intervals as their observables. Taking  $\epsilon = 0.4/H$  guarantees  $f_\pm \geq 0.6f_* > 0$ . Conditional centering preserves every  $G(\mu)$ ; the covariance projection also preserves total energy and mean flux. Number follows from the moment marginal. Nevertheless, writing  $N = \int f_* d\Gamma$ ,

$$\frac{1}{N} \int \psi^2 (f_+ - f_-) d\Gamma = \frac{2\epsilon}{N} \int f_* h^2 d\Gamma > 0. \quad (130)$$

Thus the flux populations differ even after matching the mean canonical toroidal momentum. This is stronger than a mismatch of mean momentum alone.

For mass-normalized distributions let  $\rho_\psi$  be the flux pushforward density. The conditional Jensen inequality gives the information bound

$$\frac{1}{N} \int f \ln(f/f_*) d\Gamma \geq \int \rho_\psi \ln(\rho_\psi/\rho_{\psi,*}) d\psi. \quad (131)$$

Conservation of the full flux population fixes its strictly positive right-hand side. Therefore any closed evolution satisfying (67) cannot reach  $f_*$  from  $f_+$ . Ten independent quadratures refine each

integration coordinate and all coordinates jointly. The  $32 \times 24 \times 32 \times 32$  base and  $40 \times 30 \times 40 \times 40$  joint refinement give  $\Delta\langle\psi^2\rangle = 2.5868566 \times 10^{-4}$  and a flux information floor of  $1.5140605 \times 10^{-5}$  per particle. Discrete number, energy, each moment-bin population and mean flux match to rounding; the independent squared-residual identity has about  $4 \times 10^{-12}$  relative discrepancy. These are matched *discrete* constraints; replacing sums by conditional integrals gives the continuous construction. The quadrature convergence is on the stated finite domain, not a tail limit.

The finite-domain distributions are not asserted to satisfy a confined ideal evolution. In particular,  $u = \pm 4$  walls are not tangent to the mirror force. Collision-only natural boundaries suffice for their local invariant algebra; the combined theorem separately requires closed flux-compatible ideal boundaries. Specular reversal of  $u$  at latitude endcaps preserves energy, moment and flux, but does not by itself close finite- $u$  walls. No kinetic trajectory, approach rate or physical dipole confinement time is inferred from this constructive obstruction.

## References

- [1] Naoki Sato and Philip J. Morrison. Scattering theory in noncanonical phase space: A drift-kinetic collision operator for weakly collisional plasmas. *Physics of Plasmas*, 32:102306, 2025.
- [2] Naoki Sato and Philip J. Morrison. A collision operator for describing dissipation in noncanonical phase space. *Fundamental Plasma Physics*, 10:100054, 2024.
- [3] C. Bressan, Michael Kraus, Omar Maj, and Philip J. Morrison. Metriplectic relaxation to equilibria, 2025. arXiv:2506.09787v1.
- [4] Alain J. Brizard. A guiding-center fokker–planck collision operator for nonuniform magnetic fields. *Physics of Plasmas*, 11:4429–4438, 2004.
- [5] Joshua W. Burby, Alain J. Brizard, and Hong Qin. Energetically-consistent collisional gyrokinetics. *Physics of Plasmas*, 22:100707, 2015.
- [6] Eero Hirvijoki, Alain J. Brizard, and David Pfefferlé. Differential formulation of the gyrokinetic landau operator. *Journal of Plasma Physics*, 83:595830102, 2017.
- [7] Alain J. Brizard and H. Sugama. Metriplectic bracket for guiding-center vlasov-maxwell-landau theory, 2025. arXiv:2506.22289v2.
- [8] José A. Carrillo, Jingwei Hu, Li Wang, and Jeremy Wu. A particle method for the homogeneous landau equation. *Journal of Computational Physics: X*, 7:100066, 2020.
- [9] Natalia Tronko and Alain J. Brizard. Lagrangian and hamiltonian constraints for guiding-center hamiltonian theories, 2015. Axisymmetric canonical momentum and finite-Larmor corrections.
- [10] Claudio Di Troia. Bayesian derivation of plasma equilibrium distribution function for tokamak scenarios and the associated landau collision operator, 2015. Version 1 inspected.
- [11] N. Sato, Z. Yoshida, and Y. Kawazura. Self-organization and heating by inward diffusion in magnetospheric plasmas. *Plasma and Fusion Research*, 11:2401009, 2016. arXiv version 2 inspected.
- [12] UW Plasma. Solvax: differentiable structured linear solvers in jax. <https://github.com/uwplasma/SOLVAX>, 2026. Version 0.28.2 inspected and tested 2026-10-07.

- [13] Michael Kraus and Eero Hirvijoki. Metriplectic integrators for the landau collision operator. *Physics of Plasmas*, 24:102311, 2017.
- [14] Eero Hirvijoki, Joshua W. Burby, and Michael Kraus. Energy-, momentum-, density-, and positivity-preserving spatio-temporal discretizations for the nonlinear landau collision operator with exact h-theorems, 2018. arXiv:1804.08546v1.
- [15] Eero Hirvijoki. Structure-preserving marker-particle discretizations of coulomb collisions for particle-in-cell codes. *Plasma Physics and Controlled Fusion*, 63:044003, 2021.
- [16] Filippo Zonta, Joseph V. Puszta, and Eero Hirvijoki. Multispecies structure-preserving particle discretization of the landau collision operator, 2022. arXiv:2212.00663v1.
- [17] Jie Shen and Jie Xu. Unconditionally positivity preserving and energy dissipative schemes for Poisson–Nernst–Planck equations. *Numerische Mathematik*, 148:671–697, 2021. First-order lagged mobility, implicit logarithmic chemical potential; published Section 3.1 and Theorem 3.2.
- [18] Louis Jose and Scott D. Baalrud. A generalized boltzmann kinetic theory for strongly magnetized plasmas with application to friction, 2020. arXiv:2008.06080v1.
- [19] Richard Fitzpatrick. Landau collision operator. Plasma Physics lecture notes, Eqs. (3.93)–(3.94). Accessed 7 October 2026.
- [20] Naoki Sato. Maximum entropy states of collisionless positron-electron plasma in a dipole magnetic field. *Physics of Plasmas*, 30:042503, 2023.